



With the support of the
Erasmus+ Programme
of the European Union



Erasmus Mundus Joint Master Degree "InterMaths" Interdisciplinary Mathematics

Specialisation: Decision Making and Applications to Logistics (UAB, Barcelona)

*Máster Universitario Erasmus Mundus en
"Matemáticas Interdisciplinarias"*

UNIVERSITAT AUTÒNOMA DE BARCELONA (UAB)

*Laurea Magistrale nella classe LM-44
"Mathematical Modelling"*

UNIVERSITÀ DEGLI STUDI DELL'AQUILA (UAQ)

*Master in "Interdisciplinary
Mathematics"*

TECHNISCHE UNIVERSITÄT WIEN (TUW)

Master's Thesis Title

*A Hybrid Optimization and Machine Learning Framework for Enhanced
Economic Viability Assessment (EVA) in the European Resource Adequacy
Assessment (ERAA).*

Supervisor

Prof. Susana Serna

Candidate

Mohamed Said Tawhid Ahmed

Student ID (UAQ): 302690

Academic Year 2025/2026

Acknowledgements

I would like to express my sincere gratitude to my academic supervisor, **Prof. Susana Serna**, for her invaluable guidance, critical feedback, and unwavering encouragement throughout this research. Her expertise in systems modelling and mathematical optimisation, combined with her meticulous attention to both content and presentation, has profoundly shaped this thesis.

I extend my appreciation to the **InterMaths team** for their continuous support throughout the two years of my Master’s programme, and to **Prof. Bruno Rubino** for his guidance and foundational contributions to my academic development.

I am deeply grateful to my industry mentors at **ENTSO-E** and the entire **System Development Section** for their generous support, access to institutional knowledge, and the opportunity to engage with state-of-the-art mathematical modelling tools and methodologies. Their impressive technical studies on ensuring the security of electricity supply across Europe provided the practical context that gives this research its applied relevance.

Finally, I thank my family for their unwavering support. To my mother **Nadia**, whose daily sacrifices and resilience despite her health challenges have been the foundation of my academic journey; to my father **Said**, whose guidance is priceless, my older brother **Karim**, and my twin sister **Aliaa**—their love, encouragement, and belief in me have been a constant source of strength. This work is dedicated to them.

Abstract

Accurate and computationally efficient electricity system adequacy assessment is central to European energy policy, guiding billion-euro investment decisions and underpinning the European Network of Transmission System Operators for Electricity (ENTSO-E)’s annual European Resource Adequacy Assessment (ERAA). At its core lies the Economic Viability Assessment (EVA), a multi-year iterative optimization that evaluates generation investments across diverse future scenarios. However, the EVA’s heavy reliance on computationally intensive Linear Programming (LP) models across large Monte Carlo sets makes it impractical for rapid policy analysis and real-time planning.

This thesis proposes and implements a hybrid framework that fundamentally accelerates the EVA process by integrating two innovations: (1) a machine learning surrogate model based on the Extreme Gradient Boosting (XGBoost) algorithm that replaces the full LP solve during early EVA iterations; (2) a scenario reduction technique using k-means clustering and Principal Component Analysis (PCA) to reduce a 100-year Monte Carlo set to a statistically small number of representative years.

The surrogate model is trained on data generated from a full LP simulation of the European power system implemented in Python for Power System Analysis (PyPSA), mapping system state features—including zonal demand, renewable availability, thermal capacity, and interconnection utilization—to zonal electricity prices. Results demonstrate that the hybrid surrogate-accelerated EVA achieves substantial computational speedups relative to the all-linear programming baseline, while maintaining high surrogate prediction accuracy and limiting adequacy metric errors to within acceptable tolerances of the full simulation.

The findings contribute to the academic literature on ML-assisted power system modelling and provide a methodological prototype directly applicable to future ERAA editions. Limitations, including the use of PyPSA rather than ENTSO-E’s production tool PLEXOS¹, and the simplified network representation, are discussed alongside directions for future research.

Keywords: European Resource Adequacy Assessment (ERAA), Economic Viability Assessment (EVA), Electricity markets, XGBoost, Linear Programming, PyPSA, Monte Carlo Simulation, Power System Optimization.

¹PLEXOS is a power system modeling and optimization software used globally for market analysis, capacity planning, and operational modeling.

Contents

Acknowledgements	i
Abstract	ii
List of Acronyms	viii
1 Introduction	1
1.1 Background and Motivation	1
1.2 The ENTSO-E Context	2
1.3 Research Objectives and Questions	3
1.4 Research Contributions	3
1.5 Scope and Limitations	4
1.6 Thesis Structure	4
2 Literature Review	5
2.1 ERAA and EVA	5
2.1.1 The ERAA Framework	5
2.1.2 The Role of ERAA in the European Energy Transition	5
2.1.3 The Economic Viability Assessment (EVA)	6
2.1.4 Convergence Proof Using the Banach Fixed-Point Theorem	9
2.2 Electricity Market Structure and Modeling for Economic Viability Assessment	12
2.2.1 Introduction	12
2.2.2 The Day-Ahead Market	13
2.2.3 The Intraday Market	16
2.2.4 Ancillary Services Markets	17
2.2.5 Balancing Market	19
2.2.6 Market Modeling in the EVA	21
2.2.7 Conclusions	23
2.3 ERAA Process	24
2.3.1 Introduction	24
2.3.2 Data Collection and Harmonization	25
2.3.3 Economic Viability Assessment	27
2.3.4 Adequacy Assessment	31
2.3.5 Metrics and Policy Application	33

2.3.6	Conclusions	35
3	Linear Programming for Power System Economic Dispatch	36
3.1	Linear Programming Framework	36
3.1.1	LP Formulation for Electricity Markets	36
3.1.2	PyPSA as a LP Market Simulation Platform	37
3.2	Machine Learning in Energy Systems Modelling	37
3.2.1	Surrogate Modelling for Power System Optimization	37
3.2.2	XGBoost for Tabular Energy Data	38
3.3	Scenario Reduction for Monte Carlo Simulation	38
3.3.1	The Challenge of Multi-Year Monte Carlo Scenario Sets	38
3.3.2	Clustering Algorithms for Scenario Reduction	39
3.4	Research Gaps and Positioning	39
4	Methodology, Results, and Analysis	40
4.1	Three-Zone PyPSA Network	41
4.1.1	Network Architecture	41
4.1.2	Generation Technology Mix	42
4.1.3	Load and Renewable Profiles	43
4.1.4	Transmission Network	44
4.2	Linear Programming Economic Dispatch	45
4.2.1	Purpose	45
4.2.2	Mathematical Formulation of the Economic Viability Assessment (EVA)	45
4.2.3	Interpretation	48
4.2.4	Computational Implementation: PyPSA	48
4.2.5	Annual Revenue Results	49
4.3	Multi-Year Monte Carlo Simulation and Feature Engineering	51
4.3.1	Training Data Generation	51
4.3.2	Feature Engineering	51
4.4	XGBoost Surrogate Model	52
4.4.1	Model Training and Hyperparameter Optimization	52
4.4.2	Model Performance Results	52
4.4.3	Surrogate Diagnostics	53
4.5	PCA + K-Means Weather Scenario Reduction	53
4.5.1	Motivation and Mathematical Framework	53
4.5.2	Scenario Reduction Results	54
4.6	Hybrid EVA Iteration Loop	56
4.6.1	Algorithm Design	56
4.6.2	EVA Iteration Results	56

4.7	Discussion and Limitations	58
4.7.1	Key Findings	58
4.7.2	Limitations	58
5	Conclusions	60
5.1	Original Contributions	61
5.2	Future Research Directions	63
5.2.1	Flow-Based Market Coupling Implementation	63
5.2.2	Enhanced Surrogate Generalization	63
5.2.3	Full ERAA-Scale Deployment	63
5.2.4	Uncertainty Quantification	63
5.2.5	Intraday and Balancing Market Integration	63
A	Unit Commitment and Economic Dispatch (UCED) for Adequacy	64
A.1	Purpose	64
A.2	Mathematical Formulation	64
A.2.1	Indices	64
A.2.2	Decision Variables	64
A.2.3	Objective Function	65
A.2.4	Constraints	66
A.3	Adequacy Indicators Derived	67
B	Algorithms and Pseudocode for the Hybrid ML-LP Framework	68
B.1	Master Pipeline Orchestration	68
B.2	Network Construction (Section 4.1)	68
B.3	LP Economic Dispatch via Linopy (Section 4.2.1)	70
B.4	Multi-Year Monte Carlo Training Data Generation (Section 4.3.1)	72
B.5	Feature Engineering (Section 4.3.2)	73
B.6	XGBoost Surrogate Model (Section 4.4) — Gradient-Boosted Trees	74
B.7	Scenario Reduction via PCA and K-Means Clustering (Section 4.5)	75
B.8	The Missing-Money Algorithm: Hybrid EVA Iteration Loop (Section 4.6)	76
	Bibliography	79

List of Figures

2.1	Electricity markets and their time operation.	12
2.2	Electricity trading phases.	13
2.3	Demand and Supply Curves.	13
2.4	The market clearing mechanism and Equilibrium Point.	14
2.5	Balance between electricity generation and electricity consumption. . .	14
2.6	The Day-Ahead Market.	15
2.7	Ancillary Services Markets.	18
2.8	The sequential activation of reserves.	20
2.9	ERAA process.	24
2.10	Data Sources.	26
2.11	Economic Viability Assessment.	27
2.12	Iterative process of the missing money analysis.	28
2.13	Multiyear analysis	29
2.14	Adequacy Assessment.	31
2.15	The overall ERAA Framework.	34
4.1	Network Topology.	42
4.2	Demand and Renewable Profiles (14-day excerpt).	43
4.3	Dispatch Stack (DE zone, first 168 hours).	50
4.4	Price Duration Curve.	50
4.5	Surrogate Diagnostics: MAE Comparison.	53

List of Tables

4.1	PyPSA network component summary.	41
4.2	Generator data used in the European power system model.	42
4.3	Annual generator revenues from full-year LP dispatch.	49
4.4	Full Feature Set for XGBoost Surrogate Model (DE Zone Target). . . .	51
4.5	XGBoost Hyperparameter Search Space.	52
4.6	Surrogate Model Benchmark Results on Held-Out Test Set.	52
4.7	Silhouette scores for k-means scenario reduction sweep.	54
4.8	$K = 8$ Representative Scenarios with Probability Weights.	55
4.9	EVA cost parameters.	56
4.10	Hybrid EVA iteration history.	57

List of Acronyms

EENS	Expected Energy Not Served vii, 3, 35, 64
ENTSO-E	European Network of Transmission System Operators for Electricity ii, vii, 1
ERAA	European Resource Adequacy Assessment vii, 1, 12, 24
EV	Electric Vehicle vii, 1
EVA	Economic Viability Assessment vii, 1, 12, 24
FCR	Frequency Containment Reserve vii, 18
FRR	Frequency Restoration Reserve vii, 18
LOLE	Loss of Load Expectation vii, 3, 35, 64
LP	Linear Programming ii, vii, 5, 8, 28, 36
ML	Machine Learning vii, 2
PCA	Principal Component Analysis ii, vii, 39, 40, 54
PyPSA	Python for Power System Analysis ii, vii
RES	Renewable Energy Sources vii, 1
RR	Replacement Reserves vii, 18
SRMC	Short-Run Marginal Cost vii
ST	Short-Term vii
TH	Time Horizon vii
TSO	Transmission System Operator vii, 1
TYs	Target Years vii
UC	Unit Commitment vii
UCED	Unit Commitment and Economic Dispatch vii
VoDE	Value of Dump Energy vii
VoLL	Value of Loss Load vii
VOM	Variable Operation & Maintenance vii
VoRS	Value of Reserve Shortage vii
WACC	Weighted Average Cost of Capital vii
WSs	Weather Scenarios vii
XGBoost	Extreme Gradient Boosting ii, vii, 38, 40

Chapter 1

Introduction

1.1 Background and Motivation

The European electricity sector is undergoing a profound structural transformation. The accelerated phase-out of thermal generation capacity, driven by the European Union (EU) climate targets and carbon pricing mechanisms, is creating unprecedented challenges for resource adequacy planning. At the same time, the rapid deployment of variable Renewable Energy Sources (RES), the electrification of transport through Electric Vehicles (EVs), and the emergence of distributed flexibility resources are fundamentally altering the supply-demand dynamics of continental electricity markets [1].

Against this backdrop, ensuring that sufficient generation capacity exists to meet demand reliably—defined as system adequacy—has become one of the most pressing concerns for European Transmission System Operators (TSOs) and energy regulators. An inadequate system exposes consumers to supply interruptions, drives up costs, and undermines confidence in the energy transition. Conversely, an over-built system wastes capital that could otherwise finance the transition to a net-zero energy system.

The European Resource Adequacy Assessment (ERAA), coordinated annually by the European Network of Transmission System Operators for Electricity (ENTSO-E), is the principal instrument through which the European TSOs collectively evaluate the energy system adequacy [2, 3]. At its methodological core lies the Economic Viability Assessment (EVA), a computationally intensive iterative process that simulates long-run equilibrium in generation investment, determining whether existing and planned capacity investments are economically viable under a wide range of future climate and market scenarios [2]. The EVA applies a Monte Carlo simulation framework over typically 100 or more weather years, solving a detailed LP market dispatch model for each year in each iteration, making it one of the most computationally demanding processes in European energy policy.

The computational burden of the EVA creates a fundamental tension: the analysis is rich enough to be policy-relevant, but slow enough to preclude the rapid sensitivity analyses that policymakers, investors, and TSOs urgently need. A single EVA run

can take days of high-performance computing time. When multiple policy scenarios must be compared—such as different EV penetration levels, varying carbon prices, or alternative capacity mechanisms—the computational cost becomes prohibitive.

This thesis addresses this challenge directly, proposing and implementing a hybrid computational framework that substantially accelerates the EVA process without sacrificing analytical rigour. The framework combines Machine Learning (ML) surrogate modelling, statistical scenario reduction, and EV flexibility modelling within a transparent, open-source PyPSA implementation [4].

1.2 The ENTSO-E Context

This research was conducted in close collaboration with ENTSO-E during a six-month internship placement. ENTSO-E currently uses PLEXOS as its primary production-grade market simulation tool, a commercially licensed software with powerful capabilities but limited transparency and Python integration. This thesis uses PyPSA as a research prototype environment—not to replace PLEXOS, but to demonstrate and validate the hybrid methodology in an open, reproducible framework that can inform future ERAA methodological development.

The practical motivation for this work reflects real institutional needs at ENTSO-E. As the ERAA methodology evolves towards higher spatial resolution, more Monte Carlo years, and finer temporal granularity, the computational scalability of the EVA process becomes an increasingly binding constraint. The hybrid framework developed here provides a concrete methodological pathway towards a scalable, ML-assisted EVA that could be integrated into future ERAA editions, potentially from ERAA 2026 or 2027 onwards.

1.3 Research Objectives and Questions

The overarching objective of this thesis is to design, implement, and validate a hybrid ML-LP framework for European electricity market adequacy assessment that is significantly faster than the conventional LP-only EVA, while maintaining sufficient decision-making framework for policy-relevant conclusions.

This overarching objective is decomposed into three specific research questions:

- **RQ1:** Can a machine learning surrogate model (XGBoost) accurately predict zonal electricity prices from system state features, achieving a mean absolute error below 3 €/MWh and an R^2 above 0.90 when benchmarked against full LP solutions?
- **RQ2:** How much computational speedup can the hybrid surrogate-accelerated EVA achieve relative to the all-LP baseline, and what is the corresponding trade-off in adequacy metric accuracy?
- **RQ3:** Can k-means clustering with PCA-based scenario reduction reduce the EVA's 100-year Monte Carlo set to ten representative years while maintaining adequacy metric errors within 10% of the full simulation?

1.4 Research Contributions

This thesis makes the following original contributions to the academic literature and energy industry practice:

- **Methodological contribution:** The first open-source PyPSA implementation of a surrogate-accelerated EVA framework explicitly designed to replicate and accelerate ENTSO-E's ERAA methodology.
- **Empirical contribution:** A comprehensive benchmarking study of XGBoost surrogate models for zonal price prediction in multi-zone European electricity market simulations.
- **Scenario reduction contribution:** A rigorous validation of k-means clustering for EVA scenario reduction, including sensitivity analysis of the optimal cluster number and its impact on Loss of Load Expectation (LOLE) and Expected Energy Not Served (EENS) metrics.
- **Decision-making and Policy contribution:** Concrete, quantified recommendations for incorporating ML-assisted EVA into future ERAA methodology, directly relevant to ENTSO-E's planning processes.

1.5 Scope and Limitations

This research focuses on a multi-zone representation of the European electricity system, implemented using PyPSA-Eur’s pre-clustered network data. The modelled system encompasses all synchronous zones of the ENTSO-E interconnected system, with network data sourced from publicly available datasets. The temporal scope spans a 20-year historical weather period for model training and a 100-year synthetic Monte Carlo dataset for scenario analysis.

Several important limitations should be acknowledged at the outset. First, PyPSA is used as a research prototype rather than ENTSO-E’s production tool (PLEXOS), which limits direct comparability with official ERAA results. Second, the LP dispatch model uses a linearized DC approximation (Direct Current approximation), which neglects reactive power flows and voltage constraints relevant to transmission planning but not to the adequacy assessment. These limitations are discussed in detail in Chapter 5.

1.6 Thesis Structure

The remainder of this thesis is organized to progressively introduce the theoretical background, methodological framework, and analytical results of the study as follows:

- **Chapter 2** reviews the literature on the ERAA framework, electricity market modeling, and the EVA process.
- **Chapter 3** presents the linear programming formulation, machine learning surrogate modeling, and scenario reduction techniques used in this study.
- **Chapter 4** describes the three-zone PyPSA network and reports the results of the LP dispatch, Monte Carlo simulation, XGBoost surrogate, PCA-K-means scenario reduction, and the hybrid EVA iteration loop.
- **Chapter 5** concludes the thesis, highlighting original contributions and outlining directions for future research.

Chapter 2

Literature Review

2.1 ERAA and EVA

2.1.1 The ERAA Framework

Resource adequacy assessment in interconnected electricity systems has evolved substantially over the past two decades. Early national adequacy studies, which evaluated each country’s generation surplus in isolation, proved inadequate as cross-border interconnection deepened and renewable energy penetration grew. The European Commission’s Regulation (EU) 2019/943 on the internal market for electricity mandated a harmonized European resource adequacy assessment [5], leading to the formal establishment of the ERAA as a standardized annual publication by ENTSO-E.

The ERAA methodology, as documented in ENTSO-E’s methodological guideline [6], rests on three analytical pillars. The first is the scenario framework, which defines a set of possible futures (typically a ‘central’ and ‘high RES’ scenario) characterized by generation capacity assumptions, demand projections, fuel and carbon prices, and hydro conditions. The second is the Monte Carlo weather year approach, which samples historical and synthetic climate data to generate a large ensemble of renewable generation and demand profiles, capturing inter-annual variability. The third is the market simulation, which applies a detailed LP economic dispatch model to each weather year to derive equilibrium zonal prices and generation dispatch patterns.

2.1.2 The Role of ERAA in the European Energy Transition

The ERAA serves as the cornerstone of Europe’s security of supply framework, fulfilling three critical functions that directly inform the economic viability assessment and investment signals that are the focus of this thesis. First, the ERAA assesses adequacy risks of the European power system in the medium term over the coming decade, with specific focus on target years 2028, 2030, 2033, and 2035. This forward-looking probabilistic assessment evaluates whether the power system can reliably meet demand under thousands of plausible scenarios—accounting for weather variability, forced outages, and evolving demand patterns—thereby identifying potential supply-demand

mismatches before they materialize [7]. Second, the ERAA informs decision-makers and stakeholders by serving as the reference for Member States to introduce capacity mechanisms, which are essential for ensuring that sufficient generation capacity remains available to guarantee security of supply. The Economic Viability Assessment (EVA) is central to this function, as it determines which capacities are economically viable and therefore likely to remain in or enter the market, directly influencing capacity remuneration design across Europe [8]. Third, the ERAA strengthens Europe’s energy transition through transparent data that complements other planning studies, such as the Ten-Year Network Development Plan (TYNDP). By providing a robust, data-driven foundation for understanding the economic viability of different technologies—from renewable generation to storage and demand-side flexibility—the ERAA enables the European energy transition to proceed on a solid evidence base. However, as detailed in the ERAA 2025 methodology, the current EVA process faces significant computational limitations that constrain scenario exploration and result in revenue inconsistency between the EVA and the subsequent economic dispatch adequacy model. This thesis addresses these limitations by developing a hybrid optimization and machine learning framework that enhances the EVA’s computational efficiency and accuracy, thereby enabling more robust investment signals that can better support Europe’s energy transition objectives.

2.1.3 The Economic Viability Assessment (EVA)

The EVA is the most computationally intensive component of the ERAA. Its purpose is to ensure long-run equilibrium in generation investment: the modelled capacity mix should be economically sustainable, meaning that each technology earns sufficient revenues to cover its costs, and no technology that is not present could profitably enter the market. The EVA achieves this through an iterative process: starting from an initial capacity assumption, the market model is run to compute revenues for each technology in each zone; capacities are then adjusted based on whether revenues exceed or fall short of costs; and the process repeats until convergence, defined as negligible changes in capacity between iterations [1, 9].

Mathematically, the EVA solves a fixed-point problem. For a given capacity vector K , the market simulation yields revenues $R(K)$. Capacity adjustments are made according to an entry/exit rule: capacities of technologies earning positive net revenues increase, while those earning negative net revenues decrease. Convergence to a long-run equilibrium is guaranteed under certain market structure assumptions [10], though in practice convergence can require 10–20 or more iterations, each involving a full multi-year market simulation. The total computational cost can reach hundreds of thousands of LP solves. Here is the mathematical formulation (as a fixed-point problem in the space of

capacity vectors) and explanation of the EVA Fixed-Point Problem.

Notation

Let:

- $K \in \mathbb{R}_+^N$ denote the vector of installed capacities across N technology categories or individual generation units.
- $R(K) : \mathbb{R}_+^N \rightarrow \mathbb{R}^N$ denote the revenue function, where $R_i(K)$ represents the total annual revenue (€/year) earned by technology i given capacity vector K .
- $C_i(K)$ denote the total annual cost (€/year) for technology i , including:
 - Fixed Operations and Maintenance (FOM).
 - Annualized Capital Expenditure (CAPEX).
 - Variable costs (embedded in revenues through marginal cost deduction).
- $\Pi_i(K) = R_i(K) - C_i(K)$ denote the net profit (€/year) of technology i .

The Entry/Exit Decision Rule

For each technology i , the capacity adjustment rule is defined by:

$$\Delta K_i = \Phi(\Pi_i(K)) \quad (2.1)$$

where Φ is a monotonic function satisfying:

1. $\Phi(\Pi_i) > 0$ if $\Pi_i > 0$ (positive profits $\rightarrow$ new entry)
2. $\Phi(\Pi_i) < 0$ if $\Pi_i < 0$ (negative profits $\rightarrow$ exit/decommissioning)
3. $\Phi(\Pi_i) = 0$ if $\Pi_i = 0$ (equilibrium)

In the simplest implementation, this can be expressed as:

$$K_i^{(t+1)} = K_i^{(t)} + \alpha_i \times \Pi_i(K^{(t)}) \quad (2.2)$$

where $\alpha_i > 0$ is an adjustment rate (or step size) that controls the speed of convergence.

The Fixed-Point Condition

The long-run equilibrium is achieved when the capacity vector K^* satisfies:

$$K^* = K^* + \Phi(\Pi(K^*)) \quad (2.3)$$

or equivalently:

$$\Pi_i(K^*) = 0 \quad \text{for all technologies } i \text{ with } K_i^* > 0 \quad (2.4)$$

$$\Pi_i(K^*) \leq 0 \quad \text{for all technologies } i \text{ with } K_i^* = 0 \text{ (no new entry condition)} \quad (2.5)$$

This is the standard competitive equilibrium condition: no technology earns positive profits (which would attract new entry), and no technology earns negative profits (which would cause exit).

Convergence and Computational Complexity

Convergence to the equilibrium is achieved through an iterative process:

$$K^{(0)} \rightarrow K^{(1)} \rightarrow K^{(2)} \rightarrow \dots \rightarrow K^{(M)} \approx K^{(T)} \quad (2.6)$$

where M is the number of iterations required for convergence (typically 10–20 in practice). At each iteration t , we must:

1. Solve the market simulation for $K^{(t)}$ to obtain $R(K^{(t)})$.
2. Compute $\Pi(K^{(t)}) = R(K^{(t)}) - C(K^{(t)})$.
3. Update $K^{(t+1)} = K^{(t)} + \Phi(\Pi(K^{(t)}))$.

Each market simulation involves solving a LP problem for each hour of the year across multiple weather scenarios and forced outage samples. If we denote:

- $H = 8,760$ hours per year
- $S =$ number of weather scenarios (36 full, 3 in EVA)
- $F =$ number of forced outage samples (15 full, 1 in EVA)
- $N =$ number of study zones (35 explicitly modeled)
- $G =$ number of generator units (thousands)

Then the number of LP solves per iteration is on the order of:

$$L = H \times S \times F \quad (\text{with temporal aggregation reducing this in practice}) \quad (2.7)$$

The total computational cost across all iterations is:

$$\text{Total LP solves} = M \times L \quad (2.8)$$

With $M \approx 10\text{--}20$, $S = 3$, $F = 1$, and $H = 8,760$, we get approximately:

$$\text{Total LP solves} \approx 10 \times 8,760 \times 3 \times 1 = 262,800 \text{ to } 20 \times 8,760 \times 3 \times 1 = 525,600 \quad (2.9)$$

This explains why the EVA can require hundreds of thousands of LP solves.

Convergence Guarantees

Under standard economic assumptions (convexity, continuity, and monotonicity of the excess demand function), convergence is guaranteed by the **Banach Fixed-Point Theorem** [11] if Φ is a contraction mapping:

$$\|\Phi(\Pi(K_1)) - \Phi(\Pi(K_2))\| \leq q \|K_1 - K_2\| \quad \text{for some } q \in [0, 1) \quad (2.10)$$

However, in practice with real power systems (non-convexities, discrete unit decisions, network constraints), convergence is not guaranteed and must be verified empirically.

2.1.4 Convergence Proof Using the Banach Fixed-Point Theorem

The iteration operator for the EVA is defined as:

$$T(K) = K + \alpha \Pi(K) \quad (2.11)$$

where $\alpha > 0$ is the adjustment rate.

Theorem

Let $\Pi : \mathbb{R}_+^N \rightarrow \mathbb{R}^N$ be continuously differentiable with Jacobian $J_\Pi(K)$. If $J_\Pi(K)$ is negative definite and α is chosen such that:

$$0 < \alpha < \frac{2}{\sup_K \rho(J_\Pi(K))} \quad (2.12)$$

where ρ is the spectral radius, then the iteration operator T is a contraction.

Proof

The Jacobian of T is:

$$J_T(K) = I + \alpha J_\Pi(K) \quad (2.13)$$

Since $J_\Pi(K)$ is negative definite, its eigenvalues $\lambda_i(K)$ satisfy $\lambda_i(K) < 0$. The eigenvalues of $J_T(K)$ are $1 + \alpha\lambda_i(K)$.

For T to be a contraction, we need:

$$\|J_T(K)\| = \max_i |1 + \alpha\lambda_i(K)| < 1 \quad (2.14)$$

This holds when:

$$-2 < \alpha\lambda_i(K) < 0 \quad \forall i \quad (2.15)$$

Since $\lambda_i(K) < 0$, this is equivalent to:

$$\alpha < -\frac{2}{\lambda_i(K)} \quad \forall i \quad (2.16)$$

Therefore, if:

$$\alpha < \frac{2}{\sup_K \rho(J_\Pi(K))} \quad (2.17)$$

then there exists $q \in [0, 1)$ such that:

$$\|J_T(K)\| \leq q < 1 \quad \forall K \quad (2.18)$$

By the Mean Value Theorem:

$$\|T(K_1) - T(K_2)\| \leq q\|K_1 - K_2\| \quad (2.19)$$

Thus T is a contraction mapping.

By the Banach Fixed-Point Theorem, T has a unique fixed point K^* satisfying:

$$T(K^*) = K^* \iff \Pi(K^*) = 0 \quad (2.20)$$

Moreover, for any initial $K^{(0)}$, the sequence $K^{(t+1)} = T(K^{(t)})$ converges to K^* with:

$$\|K^{(t)} - K^*\| \leq q^t \|K^{(0)} - K^*\| \rightarrow 0 \quad \text{as } t \rightarrow \infty \quad (2.21)$$

This proves the convergence of the EVA iteration to the long-run equilibrium. $\square$

2.2 Electricity Market Structure and Modeling for Economic Viability Assessment

2.2.1 Introduction

The Economic Viability Assessment (EVA) within the European Resource Adequacy Assessment (ERAA) fundamentally depends on accurately representing the electricity markets in which generation assets earn revenues. Understanding the market structure—from day-ahead trading to real-time balancing—is essential for modeling the profitability of generation, storage, and demand-side resources. This section provides a comprehensive overview of European electricity markets, with particular emphasis on the market phases that are explicitly modeled in the EVA and those that are simplified or omitted.




Figure 2.1: Electricity markets and their time operation.

As illustrated in Figure 2.1, electricity trading in Europe is organized across multiple sequential phases, each serving a distinct function in ensuring reliable and efficient power system operation.

Electricity is the fastest perishable commodity—it must be generated and consumed simultaneously. However, due to operational and market constraints, trading cannot occur in real-time. European electricity markets have transitioned to 15-minute trading periods, though the ERAA retains the former hourly resolution to maintain consistency with historical data and computational tractability. The EVA models the day-ahead market as the primary source of revenues for most generation assets, while ancillary services and balancing markets are either simplified or treated as separate post-processing steps.




Figure 2.2: Electricity trading phases.

2.2.2 The Day-Ahead Market

Market Design and Price Formation

The day-ahead market serves as the primary energy market where electricity for each hour of the next day is traded. This market establishes the reference prices that drive most generation revenues and investment decisions. The market clearing mechanism operates through a double auction process in which generators submit sell bids specifying offered quantities and associated prices, which are aggregated by ascending order into a step supply curve. Consumers or retailers submit demand bids, which are aggregated to form a downward-sloping demand curve. The intersection of these curves determines the clearing price and volume for each hour.




Figure 2.3: Demand and Supply Curves.

The clearing process fundamentally relies on the merit order—the ordering of generation units by their marginal costs. Each generation unit is characterized by its marginal cost, which represents the cost of producing one additional megawatt-hour of electricity. This cost typically includes fuel costs, variable operations and maintenance costs, and carbon emission costs. The aggregate supply curve is constructed by cumulating the capacities of all units with marginal costs below each price level. The system marginal price is then determined by the marginal cost of the last unit dispatched to meet demand.




Figure 2.4: The market clearing mechanism and Equilibrium Point.

Units with marginal costs below the system marginal price earn inframarginal rents equal to the difference between the market price and their marginal cost for each megawatt-hour produced. These rents constitute the primary revenue stream for most generation assets in energy-only markets. For peaking units with high marginal costs, revenues are concentrated in periods of high demand when their capacity is needed to maintain system balance. Conversely, baseload units with low marginal costs earn more stable but typically lower per-unit revenues.




Figure 2.5: Balance between electricity generation and electricity consumption.

Network Constraints and Market Coupling

In reality, the day-ahead market is not a single pan-European auction but a coupled set of zonal markets with transmission constraints. The ERAA models this through two primary mechanisms. The first is Flow-Based Market Coupling, which is implemented for the Core and Nordic regions [vandenbergh2016]. This approach considers the physical interdependencies of the transmission network by explicitly modeling the sensitivity of flows on critical network elements to changes in zonal net positions. The Power Transfer Distribution Factors quantify these sensitivities, enabling the market clearing algorithm to respect physical transmission limits while maximizing social welfare.

The second mechanism is the Net Transfer Capacity approach, which is applied to regions without flow-based coupling. Under this approach, bilateral exchange capacities constrain power flows between zones, but the interactions between multiple simultaneous exchanges are not explicitly modeled. While this is computationally simpler, it is less physically accurate than the flow-based approach.

An important feature of flow-based market coupling is Flow Factor Competition, which arises when two market transactions generate identical welfare but differ in their impact on scarce transmission capacity. In such cases, the transaction with the lower impact on constrained lines is prioritized. During scarcity conditions, this mechanism can lead to non-intuitive curtailment of price-taking orders, an effect that the ERAA models through the Local Matching and Curtailment Sharing mechanisms. These mechanisms ensure that curtailment is distributed equitably across bidding zones and that domestic demand is prioritized in the use of local generation.




Figure 2.6: The Day-Ahead Market.

2.2.3 The Intraday Market

Purpose and Function

The intraday market serves as a critical mechanism for adjusting electricity trading positions closer to real time, enabling market participants to respond to updated information about system conditions. Unlike the day-ahead market, which clears once per day for the following day's hours, the intraday market operates continuously or in multiple sessions up to the gate closure time—typically one hour before delivery. This market has grown substantially in importance with the increasing penetration of variable renewable energy sources, as the uncertainty associated with wind and solar generation decreases as the delivery time approaches.

The intraday market facilitates the correction of forecast errors that inevitably arise in the day-ahead time frame. For example, a wind farm operator may have committed to a certain generation level in the day-ahead market based on forecasts made twenty-four hours in advance. As the delivery hour approaches and weather forecasts become more accurate, the operator may need to adjust their position—selling surplus generation if actual wind conditions exceed forecasts, or purchasing additional energy if wind generation falls short. The intraday market provides the platform for these adjustments, allowing participants to balance their positions before the real-time balancing mechanism is activated.

Continuous Trading and Pricing Mechanisms

The intraday market operates under a continuous trading model in most European jurisdictions, in contrast to the auction-based day-ahead market. This continuous trading takes the form of an order book in which participants submit buy and sell orders that are matched when compatible prices are encountered. Orders can be submitted at any time during the trading period, and trades are executed immediately when matching orders exist. This continuous, anonymous trading environment provides maximum flexibility for participants to adjust their positions.

The pricing mechanism in the intraday market is fundamentally different from the day-ahead market. Rather than a single clearing price for each hour, the intraday market exhibits time-varying prices that reflect the evolving supply and demand balance as the delivery hour approaches. Prices can exhibit significant volatility, particularly during periods of rapid changes in forecasted renewable generation or unexpected demand shifts. This price volatility creates both risks and opportunities for market participants and can significantly affect the revenues of flexible generation assets.

The intraday market typically offers products in successive delivery periods, with prod-

ucts becoming available for trading at different times before delivery. While the ERAA methodology retains hourly resolution, many European intraday markets now operate on fifteen-minute or even five-minute products, enabling finer temporal granularity in position adjustments.

Relationship to the EVA

The intraday market presents both opportunities and challenges for the Economic Viability Assessment. On one hand, excluding intraday revenues may underestimate the total revenues of flexible assets, particularly those that can adjust their positions in response to intraday price signals. Storage assets, flexible gas turbines, and demand response resources can capture significant value in the intraday market by providing the flexibility needed to accommodate forecast errors and unexpected system conditions.

On the other hand, including intraday market dynamics in the EVA would dramatically increase computational complexity. The continuous trading nature of the intraday market, combined with its dependence on the evolution of forecasts and system conditions, would require fundamentally different modeling approaches than the deterministic day-ahead optimization currently employed. Moreover, the historical data required to calibrate intraday price models is often less complete and less standardized across jurisdictions than day-ahead market data.

The ERAA methodology currently addresses this challenge by treating intraday market activities as either exogenous or implicit in the day-ahead model through the assumption of perfect foresight. Under perfect foresight, all relevant information is assumed to be known in advance, eliminating the forecasting errors that drive intraday trading. While this simplification is necessary for computational tractability, it represents a limitation of the current approach that future methodological developments may address.

2.2.4 Ancillary Services Markets

Classification and Purpose

Ancillary services are essential for maintaining power system security and reliability, encompassing a range of functions including frequency control, voltage support, and system restoration. The ERAA focuses exclusively on frequency-related services—collectively referred to as reserves—as these directly impact system balance and adequacy assessments. Three main reserve categories are modeled within the ERAA framework.

Electricity markets cannot rely solely on the day-ahead energy market because the




Figure 2.7: Ancillary Services Markets.

power system requires ancillary services to maintain frequency, voltage, and operational security. The main types of ancillary services—as shown in Figure 2.7—include:

- **Frequency Containment Reserve (FCR):** automatic response within seconds to arrest frequency deviations.
- **Frequency Restoration Reserve (FRR):** automatic or manual activation to restore frequency to nominal.
- **Replacement Reserves (RR):** slower reserves to relieve FRR.

Reserve Capacity Markets

In the ERAA context, the ancillary service market is reduced to a reserve capacity market in which generating units offer megawatts of available capacity for both upward and downward regulation. The market clearing for reserve capacity involves a procurement optimization that selects the least-cost combination of reserves subject to technical requirements and system constraints. A fundamental characteristic of reserve provision is capacity exclusivity. Capacity allocated to upward reserve provision cannot simultaneously be used to sell energy in the day-ahead market. This creates an opportunity cost that must be considered in economic viability calculations. When a unit provides upward reserves, it must forgo the energy market revenues it could have earned on the withheld capacity. Conversely, providing downward regulation requires a unit to be online and generating at least the amount of capacity offered, which may force the unit to operate at a minimum stable level even when energy market prices are low.

European Integration

European balancing processes are increasingly coordinated through dedicated platforms that harmonize balancing markets and facilitate cross-border activation of reserves. PICASSO, the Platform for the International Coordination of Automated Frequency

Restoration, coordinates aFRR across participating TSOs, enabling cross-border activation of automated reserves when this is economically efficient. MARI, the Manually Activated Reserve Initiative, performs a similar function for mFRR across Europe. These platforms represent significant steps toward a fully integrated European balancing market, improving the efficiency of reserve procurement and reducing the cost of maintaining system security.

2.2.5 Balancing Market

Activation Sequence

TSOs hold the final responsibility for maintaining the grid frequency at its nominal value at all times. When an imbalance occurs—such as a generator tripping, a forecast error in wind or solar generation, or unexpected demand changes—the TSO activates the reserved capacity to provide real-time balancing energy. The activation sequence follows the classical pattern: Frequency Containment Reserve is activated first, stabilizing frequency deviations within seconds. Once frequency is contained, Automatic Frequency Restoration Reserve is activated to restore frequency to its nominal value and bring cross-border exchanges back to their scheduled values. Finally, Manual Frequency Restoration Reserve is activated on a unit-specific basis, allowing the TSO to fine-tune the system balance and rebuild reserve margins.

This reactive balancing approach, in which balancing energy is delivered only after an imbalance has actually occurred, represents the traditional operational paradigm. However, a proactive balancing approach is increasingly used in which TSOs activate reserves to mitigate forecasted imbalances before they materialize. In this case, the activation sequence typically reverses to prioritize slower and more economical products first, reserving the faster products for real-time contingencies.

Pricing Mechanism

To ensure efficient price signals and incentivize participation, the real-time balancing market operates under marginal pricing. A single down-regulation price is defined as the lowest accepted price among the activated down-regulation bids, while the up-regulation price is set equal to the highest accepted price among the activated up-regulation bids. All activated bids are settled at the corresponding marginal regulation price, regardless of their individual bid prices. This uniform pricing mechanism provides strong incentives for truthful bidding and efficient resource allocation. The providers of reserve receive payment not only for the capacity being available but also for the delivery or curtailment of energy. The activated balancing energy is traded and priced according to the design of the balancing market, commonly using marginal pricing

mechanisms. The total remuneration for a reserve provider thus comprises a capacity payment and an energy payment, each determined through separate but interconnected market mechanisms.

The Sequential Activation of Reserves

The balancing market is the real-time mechanism that activates these reserves sequentially. Figure 2.8 illustrates how FCR acts first (seconds), then FRR (minutes), and finally RR (15 minutes to hours). In the ERAA's EVA and UCED models, these services are aggregated into a single balancing reserve requirement BR_n per zone. A portion of the available capacity must be withheld from the energy market to provide these reserves (see the reserve withholding constraint in the UCED formulation). This simplification is necessary for long-term adequacy studies but loses the temporal granularity of sequential activation.




Figure 2.8: The sequential activation of reserves.

2.2.6 Market Modeling in the EVA

Primary Focus on Day-Ahead Market

The EVA primarily models the day-ahead market as the source of revenues for generation assets. This focus is justified by the primacy of day-ahead market revenues for most generation assets, particularly baseload and mid-merit units for which day-ahead transactions constitute the majority of total revenues. Additionally, day-ahead prices drive investment signals and are the most transparent and predictable component of total revenues, making them the most suitable basis for economic viability assessment. Finally, modeling the day-ahead market at hourly resolution across multiple years and scenarios is already computationally intensive; adding intraday and balancing markets would make the problem intractable given current computational resources.

Simplified Modeling of Ancillary Services

Ancillary services are modeled in a simplified manner within the EVA. Reserve requirements are primarily modeled through capacity withholding, in which an amount of generation capacity equal to the total FCR and aFRR requirement is withheld from the day-ahead market. This approach reflects the capacity exclusivity constraint without requiring explicit modeling of reserve market clearing. The opportunity cost of providing reserves is implicitly captured through reduced day-ahead revenues. By withholding capacity from the day-ahead market, units forego potential energy market revenues, and this foregone revenue represents the opportunity cost of reserve provision. TSOs can choose to account for reserve requirements from specific technology categories—thermal, renewable, DSR, battery, and hydro—based on reported data and technical capabilities.

Balancing Market Simplification

The balancing market is not explicitly modeled in the EVA for several fundamental reasons. The perfect foresight assumption eliminates the forecasting uncertainty that drives balancing market activity, removing the primary motivation for real-time adjustments. Additionally, balancing market revenues are relatively small compared to day-ahead revenues, particularly for baseload and mid-merit units. The primary focus of the EVA is structural adequacy—whether the system is balanced in expectation—rather than real-time operational issues, which are the subject of separate operational assessments.

Revenue Consistency Between Market Phases

A critical challenge in the EVA is ensuring consistency between modeled day-ahead revenues and the revenues that units would actually earn across all market phases. The ERAA 2025 methodology acknowledges that significant inconsistencies exist between EVA and Economic Dispatch (ED) revenues, stemming from multiple sources. Scenario reduction from thirty-six to three weather years truncates the revenue distribution, while forced outage reduction from fifteen to one underestimates tail risk. Unit aggregation loses unit-specific economics, and the simplification of ancillary services undercounts total revenues. The revenue-based EVA approach directly addresses this consistency gap by using revenues from the full economic dispatch model rather than from the cost-based optimization.

The Intraday Market in Context

The intraday market has become increasingly important with the growing penetration of variable renewable energy sources. As wind and solar generation constitute larger shares of the generation mix, the uncertainty associated with forecasting their output becomes more consequential. The day-ahead time frame, with its twenty-four hour lead time, necessarily involves significant forecast uncertainty that can only be resolved as the delivery hour approaches. The intraday market provides the mechanism for resolving this uncertainty and enabling efficient adjustments to generation and consumption plans. The value of intraday trading is particularly significant for flexible resources. Storage assets, flexible gas turbines, and demand response can capture value by adjusting their positions in response to evolving price signals. Conversely, inflexible resources may be exposed to the risk of adverse price movements if they must correct forecast errors in the intraday market. Understanding these dynamics is increasingly important for accurate economic viability assessment.

Modeling Challenges

Modeling the intraday market within the EVA presents significant challenges. The continuous trading nature of the market, combined with its dependence on the evolution of forecasts and system conditions, would require fundamentally different modeling approaches than the deterministic day-ahead optimization currently employed. The absence of standardized and complete historical data across all European jurisdictions further complicates the modeling effort.

Moreover, the intraday market's role changes over time. As the delivery hour approaches, the intraday market serves different functions—providing a platform for forecast error correction in the hours immediately following the day-ahead closure, then

facilitating the fine-tuning of positions as real-time approaches, and finally serving as a last resort for balancing before the activation of reserves.

Implications for the EVA

The growing importance of the intraday market suggests that future methodological developments should consider its inclusion in the EVA. Potential approaches include exogenous representation of intraday revenues based on historical patterns, simplified intraday models that capture the key price dynamics without full continuous trading simulation, and machine learning methods that can predict intraday prices based on system conditions. Each approach presents different trade-offs between accuracy and computational tractability, and the optimal approach may depend on the specific technology and market context being assessed.

2.2.7 Conclusions

This section has provided a comprehensive overview of European electricity markets and their representation in the Economic Viability Assessment. The day-ahead market serves as the primary source of generation revenues and the foundation for the EVA's market clearing model. Ancillary services are modeled through simplified capacity withholding mechanisms that reflect capacity exclusivity and opportunity costs. The balancing market is not explicitly modeled due to the perfect foresight assumption and the focus on structural adequacy.

The intraday market, which has grown substantially in importance with the increasing penetration of variable renewable energy, presents both opportunities and challenges for the EVA. While excluding intraday revenues may underestimate the total revenues of flexible assets, including intraday dynamics would dramatically increase computational complexity. The revenue-based EVA approach offers a promising path forward by using revenues from the full economic dispatch model to calibrate the economic viability assessment.

The understanding of electricity market structure developed in this section provides the foundation for the hybrid optimization and machine learning framework proposed in this thesis. By accurately modeling the markets in which generation assets earn revenues, we can better assess their economic viability and provide more robust investment signals for the European energy transition. The computational challenges associated with intraday and balancing market modeling suggest that machine learning approaches may offer a path forward, enabling the representation of these markets without the prohibitive computational cost that would accompany full simulation of their complex dynamics.

2.3 ERAA Process

2.3.1 Introduction

The European Resource Adequacy Assessment (ERAA) represents the most comprehensive probabilistic adequacy study conducted at the pan-European level, serving as the cornerstone for security of supply assessment across the continent. This section provides a detailed examination of the ERAA process, with particular emphasis on the Economic Viability Assessment (EVA)—the critical stage that determines which generation, storage, and demand-side resources are economically viable over the 10-year planning horizon and therefore likely to enter or exit the market. The EVA stage is the central focus of this thesis, and understanding its placement within the broader ERAA workflow is essential for appreciating the methodological contributions proposed herein.

The ERAA methodology follows a structured multi-stage process that begins with extensive data collection and culminates in the publication of adequacy metrics that inform capacity mechanism design and investment decisions across Europe. This process is depicted schematically in Figure 2.9, which illustrates the sequential progression from data preparation through economic viability assessment to adequacy simulation and final metric calculation. Each stage builds upon the outputs of the previous stage, creating an integrated assessment framework that captures the complex interactions between investment decisions, market outcomes, and system reliability.




Figure 2.9: ERAA process.

The flowchart reveals the central position of the EVA within the overall workflow. Following the collection and harmonization of input data from Transmission System Operators across Europe, the process diverges into two parallel streams: the EVA

determines the economically viable capacity mix, while the subsequent Adequacy Assessment evaluates the reliability of that mix under a comprehensive set of probabilistic scenarios. The Metrics stage then translates the simulation outputs into policy-relevant indicators that can be compared against established European reliability standards.

2.3.2 Data Collection and Harmonization

Overview of Data Sources

The ERAA process begins with an extensive data collection effort that draws upon multiple sources of information about the European power system. Transmission System Operators provide detailed data on generation units, including their technical characteristics, economic parameters, and expected decommissioning dates. This information is collected through the Pan-European Market Modelling Data Base (PEMMDB), which serves as the single source of supply-side and grid data across all ENTSO-E assessments [12, 13]. The database contains unit-by-unit resolution of European power plants, their technical and economic parameters, and forecasted development of renewable energy capacities [14].

Climate-dependent data is sourced from the Pan-European Climate Database (PECD), which provides high-resolution weather information essential for modeling renewable generation and demand [15, 16]. Version 4.2 of the PECD, used for ERAA 2025, contains both historical reanalysis data from 1950 to 2024 and projected data through 2100 based on six climate models under various greenhouse gas emission scenarios [15]. The dataset includes hourly temperature, wind speed, solar irradiance, and precipitation information, as well as derived energy variables including wind power onshore and offshore, solar photovoltaic generation, and hydropower inflow and generation profiles [14].

Demand forecasting is performed using the Demand Forecasting Toolbox (DFT), a specialized software application that creates hourly demand profiles based on historical load data, weather variables, calendar effects, and technology diffusion assumptions [17]. The DFT incorporates uncertainty analysis under various climate conditions and can be customized to reflect the specific characteristics of different market nodes. Some TSOs choose to provide their own demand time series using their own internal forecasting tools, which are then incorporated into the ERAA data pipeline [17].

Data Processing and Scenario Generation

The raw input data undergoes extensive processing to prepare it for the subsequent modeling stages. Generation data from the PEMMDB is categorized by technology type and aggregation level, with renewable energy sources aggregated at the PECD zone level and thermal generation maintained at unit-by-unit granularity. Weather-dependent time series for wind, solar, and hydro generation are generated from the PECD data and assigned to the appropriate generation units or aggregated capacities.

The processed data is organized into the National Trends scenario, which represents the central projection of power system evolution over the planning horizon. The National Trends scenario incorporates TSOs' best estimates of future generation capacity, demand, and network development, aligned with national development plans and policy targets. This scenario serves as the baseline for both the EVA and the subsequent adequacy assessment.

Beyond the National Trends scenario, additional scenarios are generated to capture uncertainty in key drivers such as fuel prices, carbon prices, and technology costs. These scenarios are used to test the sensitivity of adequacy results to alternative assumptions and to provide a range of outcomes that inform decision-making under uncertainty.




Figure 2.10: Data Sources.

2.3.3 Economic Viability Assessment

Purpose and Scope

The Economic Viability Assessment is the stage of the ERAA process that determines which generation, storage, and demand-side resources are economically viable over the 10-year planning horizon. This assessment is essential because the capacity mix used in the adequacy assessment must reflect realistic market outcomes rather than assuming that all existing capacity remains available and all potential new capacity is built. Without the EVA, the adequacy assessment would potentially overestimate system reliability by including capacity that would have been retired due to economic pressures.

The EVA scope includes all generation units that primarily depend on electricity-only market revenues. For these units, the EVA determines whether they will remain in operation, be decommissioned, be mothballed or de-mothballed, or have their lifetimes extended. For potential new capacity, the EVA determines which investments are economically viable and therefore likely to enter the market. Technologies covered by the EVA include natural gas units (eligible for all decisions), lignite and hard coal units, oil-fired units, hydrogen-fired turbines, demand-side response, and battery storage with a six-hour energy-to-power ratio.




Figure 2.11: Economic Viability Assessment.

The Iterative EVA Process

The EVA operates through an iterative process that alternates between market simulation and capacity adjustment. This iterative nature is essential because the economic viability of any individual unit depends on the overall capacity mix, which in turn depends on the viability of all other units. The process begins with the National Trends

capacity mix for the first target year and proceeds through a sequence of simulations and capacity adjustments until convergence is achieved.

Each iteration of the EVA involves solving a LP-based economic dispatch model that determines the optimal operation of the power system for the given capacity mix. This simulation covers all hours of the year for a reduced set of weather scenarios and forced outage samples, capturing the key sources of uncertainty while maintaining computational tractability. The simulation produces zonal electricity prices that reflect the marginal cost of the last dispatched unit in each hour.




Figure 2.12: Iterative process of the missing money analysis.

From these prices, revenues are calculated for each generation unit, and these revenues are compared against the unit's costs. Units with positive net revenues are deemed economically viable and retained in the capacity mix, while units with negative net revenues are considered for decommissioning, mothballing, or lifetime extension. New entry decisions are evaluated by assessing the viability of potential new units given the current capacity mix and market conditions.

The capacity adjustments made at each iteration are typically modest—the process does not remove all unviable units in a single step but rather removes the most unviable units progressively. This gradual approach helps maintain numerical stability and prevents overshooting, where the removal of too many units in one iteration would create scarcity rents that would make previously unviable units profitable again. The process continues until all remaining units are profitable and no additional new entry would be profitable, representing a long-run equilibrium.

Weather Scenario Reduction

A critical feature of the EVA is the reduction of weather scenarios from the full set of thirty-six to a much smaller subset of typically three. This reduction is necessary because the iterative EVA process would be computationally prohibitive if the full scenario set were used. The selected subset must be as representative as possible of the full set to ensure consistency between the EVA and the subsequent adequacy assessment, which uses all thirty-six scenarios.




Figure 2.13: Multiyear analysis

The scenario selection process aims to minimize the differences in the distribution of net revenues from thermal technologies between the subset and the complete set of weather scenarios. The Wasserstein distance, which measures the area between the empirical cumulative distribution functions of revenues, is used as the optimization objective. The selected subset is the one that minimizes the total Wasserstein distance across all study zones.

The ERAA 2025 scenario selection was based on the previous year's adequacy results, which improved the expected representativeness and consistency with respect to the current year's results. The selection process accounted for all target years and forced outage samples to further improve robustness. Unlike the previous year's approach, which assigned weights to the selected scenarios, the current approach aims to identify the most representative subset of equally weighted scenarios, with no explicit weights applied.

Unit Aggregation and Result Transfer

To keep the computational complexity of the EVA tractable, generator units are aggregated at the nodal level based on their main characteristics: technology, fuel, and techno-economic parameters. This aggregation is possible because the EVA model

solves in a linear manner and uses a uniform derating of net generating capacities based on forced outage rates rather than random draws of outage patterns.

The aggregation of units, while necessary for computational tractability, creates a challenge for the transfer of EVA results to the subsequent adequacy assessment. The adequacy model operates at unit-by-unit granularity, requiring a mapping from the aggregated EVA outcomes to individual units. The ERAA 2025 methodology addresses this challenge through a unit-based decommissioning approach: within each aggregated EVA unit, units are processed in a deterministic order (oldest commissioning date, then earliest decommissioning year, then alphabetical) and derated sequentially until the EVA capacity change is reached.

This unit-based approach represents a significant improvement over the previous linear derating approach, which applied capacity reductions homogeneously across all units in a technology class. The new approach produces more realistic inputs for the adequacy assessment: units are removed as whole assets where possible, and at most one unit is partially derated to close any remaining capacity gap. This avoids spreading small residuals of capacity across many units and yields outage and availability patterns that better reflect an actual fleet, while respecting the EVA's objective of determining overall viable capacity per technology in each study zone.

2.3.4 Adequacy Assessment

Monte Carlo Simulation Framework

Following the EVA, the adequacy assessment evaluates the reliability of the economically viable capacity mix through a comprehensive Monte Carlo simulation. This simulation captures the key uncertainties affecting system adequacy—weather variability and forced outages—through a large number of scenario realizations. The combination of multiple weather scenarios and multiple forced outage samples yields a comprehensive set of possible system states that can be assessed probabilistically.

The Monte Carlo simulation begins by defining the weather scenarios that will be used. Each weather scenario represents a consistent set of temperature-dependent demand time series, wind and solar load factor time series, hydro generation and inflow time series, and climate-dependent time series for other renewable and non-renewable generation. These weather scenarios are drawn from the PECD and represent different possible evolutions of climate over the planning horizon.

For each weather scenario, multiple sets of random forced outage realizations are generated. These outages affect thermal generation units and transmission assets (HVDC and HVAC interconnections), with the impact on installed capacities known during the Unit Commitment and Economic Dispatch step. The forced outage profiles are generated randomly based on forced outage rates and mean time to repair parameters provided by TSOs or drawn from default values.




Figure 2.14: Adequacy Assessment.

Unit Commitment and Economic Dispatch

The Unit Commitment and Economic Dispatch model represents the core simulation engine of the adequacy assessment. This two-step optimization minimizes system costs while respecting the physical and operational constraints of the power system. The first step involves an annual optimization that accounts for intertemporal energy constraints that might span the whole year, such as hydro reservoir management and seasonal storage operation. Multiple hours are aggregated and optimized in blocks to maintain computational tractability.

The outcome of the annual optimization comprises granular daily target values for objects with annual constraints, particularly hydro reservoirs. These daily reservoir targets are set as soft boundaries to the total hydro energy available over the day for the subsequent, more granular optimization step. The UCED optimization is then performed in smaller time steps, typically one day, to determine which units are dispatched for each hour of the optimization horizon and the respective dispatch level for each unit.

The UCED problem is solved for each hour of the year, with the system state defined by demand, renewable feed-in, available thermal generation, and network constraints. The solver employs flexible hydro storage resources such as reservoirs and pumped storage plants to exploit marginal price gain opportunities from a cost minimization perspective. Final marginal prices are set equal to the highest marginal cost of the dispatched resources to cover the hourly domestic demand, reflecting the marginal cost of the last unit dispatched.

Convergence and Adequacy Indicators

The Monte Carlo simulation continues until convergence is achieved in the adequacy indicators. Convergence is assessed using the relative change of the coefficient of variation derived from the Expected Energy Not Served of the entire geographical scope. The coefficient of variation measures the standard deviation of the EENS estimate relative to its mean, providing an indication of the stability of the results. No explicit simulation stopping criterion is set for the coefficient of variation; instead, the decision whether to launch additional model runs is based on a compromise between the relative change in the coefficient and the required computational time.

The key adequacy indicators calculated from the simulation are Loss of Load Expectation and Expected Energy Not Served. LOLE measures the expected number of hours during which resources are insufficient to meet demand, expressed in hours per year. This indicator captures the frequency of inadequacy events but does not indicate their severity. EENS measures the electricity demand that is expected not to be supplied

due to insufficient resources, expressed in gigawatt-hours per year. This indicator captures both the frequency and severity of inadequacy events, providing a more complete picture of reliability risk.

The final adequacy indicators reflect the impact of the Curtailment Sharing implementation in the adequacy assessment. Curtailment Sharing is a component of the EUPHEMIA market coupling algorithm that ensures the equitable distribution of curtailment across bidding zones during scarcity conditions. The ERAA implementation of Curtailment Sharing uses Domestic Energy Not Served as a proxy for price-taking order volume, with a quadratic penalty function that equalizes curtailment ratios across zones subject to network constraints.

2.3.5 Metrics and Policy Application

Interpretation of Results

The LOLE and EENS metrics produced by the adequacy assessment serve as the primary indicators of security of supply risk. These metrics are reported at multiple levels of aggregation: pan-European, by country, and by study zone. The pan-European metrics provide a high-level assessment of overall system reliability, while the country-level and zone-level metrics enable identification of specific regions that may face adequacy challenges.

The metrics are compared against European reliability standards to determine whether the system meets the established security of supply criteria. The most commonly referenced standard is LOLE less than three hours per year, which represents the target reliability level considered acceptable for European power systems. EENS is typically considered alongside LOLE to provide additional insight into the severity of expected shortfalls.

The results are presented as a range rather than a single point estimate to reflect the uncertainty inherent in the assessment. The range arises from the two risk-aversion modeling approaches implemented in the EVA: the lower bound corresponds to the hurdle premium only approach, while the upper bound corresponds to the hurdle premium combined with revenue cap approach. This range acknowledges that investor behavior under uncertainty cannot be precisely predicted and that different assumptions about risk aversion can lead to different adequacy outcomes.

Decision-Making and Policy Implications

The ERAA results have significant policy implications at both the European and national levels. At the European level, the results inform the European Commission's assessment of security of supply across the continent and support the development of coordinated policies to address adequacy challenges. The results also provide input to the Ten-Year Network Development Plan and other system planning studies.

At the national level, the ERAA results serve as the reference for Member States to introduce capacity mechanisms. The Electricity Regulation requires that Member States with identified adequacy risks implement measures to ensure security of supply, which may include capacity remuneration mechanisms, strategic reserves, or other market interventions. The ERAA provides the evidence base for these decisions and ensures that capacity mechanisms are justified and appropriately designed.

The transparency of the ERAA process and the availability of detailed results enable stakeholders to understand the drivers of adequacy risks and the potential impacts of alternative policies. The stakeholder consultation process, which includes bilateral exchanges and public consultations, ensures that diverse perspectives are considered in the methodology development and that the results are credible and widely accepted.




Figure 2.15: The overall ERAA Framework.

2.3.6 Conclusions

The ERAA follows a structured probabilistic process, as depicted in the flowchart. The key stages are:

1. **Data:** Transmission system operators (TSOs) share data on generation, demand, grid topology, and weather-dependent renewable profiles.
2. **Economic Viability Assessment (EVA)** – the focus of this thesis. EVA determines which existing generation units remain economically viable and which new units (gas, batteries, hydrogen, DSR) would be built over a 10-year horizon. It uses a cost-minimizing linear program (LP) under a limited set of weather scenarios (typically 3 out of 36).
3. **Adequacy:** A full Monte Carlo simulation over many weather years (36) and forced outage patterns (15) using a Unit Commitment & Economic Dispatch (UCED) model. This stage computes the core adequacy metrics: LOLE in hours/year and EENS in GWh/year.
4. **Metrics:** LOLE and EENS are compared against European reliability standards (e.g., $\text{LOLE} < 3 \text{ hours/year}$).

The EVA step is iterative: investment decisions affect dispatch, and resulting revenues feed back into entry/exit decisions. This iteration is computationally heavy, motivating the methodological developments described later.

Chapter 3

Linear Programming for Power System Economic Dispatch

3.1 Linear Programming Framework

3.1.1 LP Formulation for Electricity Markets

The economic dispatch problem—determining the least-cost allocation of generation to meet demand subject to network and technical constraints—is a foundational problem in power systems engineering. In its simplest form, the economic dispatch minimizes total generation cost subject to power balance, generation capacity, and transmission constraints. This is inherently a linear programme when generation costs are linearized (i.e., marginal costs are constant or piecewise-linear) and transmission is modelled using a DC approximation.¹

The DC optimal power flow (DC-OPF) model, widely used in academic and commercial adequacy studies, represents transmission as a linear flow model where power flows are proportional to phase angle differences. Under this approximation, the full nonlinear AC (Alternating Current) power flow equations are linearized, yielding a tractable LP that can be solved efficiently even for large multi-zone systems. The trade-off is the neglect of reactive power and voltage constraints, which are not material for adequacy assessment but are critical for transmission planning.

The mathematical foundation of LP traces back to the seminal work of Dantzig [18], who developed the simplex method for solving linear programming problems. In the context of power systems, Schweppe [19] established the theoretical framework for spot pricing of electricity, demonstrating that marginal prices emerge naturally as the dual variables of the power balance constraints in the economic dispatch problem. This

¹Under the DC approximation, transmission losses are neglected, voltage magnitudes are assumed to be Unity, and reactive power is ignored. This yields a linear relationship between active power flows and bus voltage angles, reducing the AC power flow equations to a set of linear constraints that allow efficient solution via linear programming. While this approximation is widely accepted for adequacy studies and market simulations, it does not capture voltage stability, reactive power dispatch, or congestion due to reactive power limits, which is beyond the scope of this thesis.

insight underpins the revenue calculation methodology used in the EVA.

3.1.2 PyPSA as a LP Market Simulation Platform

PyPSA (Python for Power System Analysis) is an open-source Python package for simulating and optimizing modern power systems, developed primarily by [4] and colleagues at KIT and TU Berlin. PyPSA implements a full LP and MILP optimization framework for multi-zone electricity systems, including optimal power flow, storage dispatch, and unit commitment. Its key advantages for this research include: (1) full Python integration, enabling seamless data pipelines and ML model interfacing; (2) transparency and reproducibility, as all model equations are accessible and documented; (3) a growing ecosystem, including PyPSA-Eur, a pre-configured European network model; and (4) solver flexibility, supporting both commercial solvers (Gurobi, CPLEX) and open-source alternatives (HiGHS, GLPK).

PyPSA-Eur, provides the network data infrastructure for this research, offering pre-clustered European networks with varying spatial resolution (from 37 to 1000+ nodes), historical renewable generation profiles derived from ERA5² reanalysis data, and demand profiles calibrated to ENTSO-E’s transparency data³. This infrastructure substantially reduces the data engineering burden and allows the research to focus on methodological innovation.

3.2 Machine Learning in Energy Systems Modelling

3.2.1 Surrogate Modelling for Power System Optimization

Surrogate modelling—the use of fast approximate models to replace computationally expensive simulations—has a long history in engineering design optimization. In the context of power systems, surrogate models have been applied to a variety of problems including contingency analysis, security-constrained dispatch, and system planning. The fundamental premise is that a trained ML model can map inputs (system state variables) to outputs (prices, dispatch decisions) with sufficient accuracy at a fraction of the computational cost of the original LP or MILP.

The theoretical justification for surrogate modelling in power systems rests on the observation that the LP solution—and particularly its dual variables (prices)—is a deterministic function of the problem parameters (loads, capacities, fuel prices). This

²ERA5 is a global climate reanalysis dataset produced by the ECMWF through the Copernicus Climate Change Service.

³The ENTSO-E Transparency Platform is the primary public repository for European electricity market data, mandated by EU Regulation 543/2013 to publish standardized information on generation, load, transmission, and balancing.

function is piecewise-linear and continuous, but the number of possible optimal bases grows combinatorially with system size. Machine learning methods, particularly tree-based ensemble methods, can effectively approximate this complex function without explicit enumeration of the basis space.

3.2.2 XGBoost for Tabular Energy Data

Among the wide range of ML algorithms available, gradient boosting methods—and XGBoost in particular—have consistently demonstrated superior performance on tabular data compared to both linear models and deep neural networks [20, 21]. XGBoost has been developed as a scalable tree boosting system that builds an ensemble of decision trees sequentially, each correcting the errors of the previous, using a regularized objective function that prevents overfitting [20]. For the electricity price prediction task, the features are structured, tabular, and moderately high-dimensional, making XGBoost a natural choice.

Key hyperparameters governing XGBoost performance include the number of estimators (`n_estimators`), the tree depth (`max_depth`), the learning rate (η), and the subsampling ratio (`subsample`). Hyperparameter optimization via cross-validated grid search or Bayesian optimization is standard practice. The built-in feature importance mechanism of XGBoost, which ranks inputs by their contribution to prediction accuracy, provides interpretability that is particularly valuable for academic and policy audiences.

3.3 Scenario Reduction for Monte Carlo Simulation

3.3.1 The Challenge of Multi-Year Monte Carlo Scenario Sets

The ERAA’s Monte Carlo approach to capturing climate variability requires running the market simulation over a large number of weather years—typically 100 or more—to obtain statistically robust adequacy metrics. Each additional year adds linearly to the computational cost, creating a direct tension between statistical accuracy and computational tractability. Scenario reduction—selecting a small, representative subset of scenarios that preserves the statistical properties of the full set—is a standard technique in stochastic programming and has been applied to electricity market problems.

The mathematical framework for scenario reduction builds on established algorithms for reducing the number of scenarios in stochastic programming problems. The objective is to find a subset S of the original scenario set Ω such that the probability-weighted distance between the distribution of outcomes under S and under Ω is minimized. In

practice, clustering algorithms are applied to the scenario feature space (e.g., annual renewable profiles and demand patterns), and the centroid of each cluster serves as the representative scenario, weighted by the cluster size divided by the total number of scenarios.

3.3.2 Clustering Algorithms for Scenario Reduction

K-means clustering is the most widely used scenario reduction technique in energy system modelling due to its computational efficiency and simplicity [22]. In the energy context, each scenario (a full year of hourly data) is represented as a high-dimensional vector, and k-means finds k centroids that minimize the total within-cluster variance. The challenge of high dimensionality is addressed by first applying PCA to project the scenarios onto a lower-dimensional subspace that retains the dominant sources of variability [23].

The combination of PCA and k-means has been widely applied in power system analysis. Unsupervised machine learning algorithms, including principal component analysis (PCA) and K-means clustering, have proven instrumental for scenario reduction in renewable energy systems [24]. Similarly, PCA combined with clustering has been used to extract typical operating scenarios for high-penetration renewable energy power systems, helping to reduce dimensionality while preserving the statistical characteristics of the original data [25].

3.4 Research Gaps and Positioning

The review above identifies two primary gaps in the existing literature that this thesis addresses. First, while ML surrogates for power system applications are well-established, no prior work has applied surrogate modelling specifically within the EVA iterative loop, addressing the challenge of maintaining accuracy as the capacity mix evolves. Second, while k-means scenario reduction has been applied in energy system planning, it has not been rigorously validated against the ERAA’s adequacy metrics (LOLE and EENS) using a PyPSA-based implementation.

This thesis fills both gaps within a single, coherent hybrid framework implemented in open-source software, making the methodology reproducible and extensible by the research community and ENTSO-E alike.

Chapter 4

Methodology, Results, and Analysis

This chapter documents the complete computational implementation of the Hybrid Machine Learning – Linear Programming (ML–LP) framework for ERAA, developed and validated in PyPSA with the Linopy optimization interface. The implementation spans five tightly integrated modules: (1) a three-zone PyPSA network representing Germany, France, and the Netherlands; (2) a linear programming economic dispatch model solved at hourly resolution using the HiGHS open-source solver; (3) a Monte Carlo multi-year simulation framework generating labelled training data; (4) an XGBoost gradient-boosted surrogate model trained to predict zonal marginal prices; and (5) a PCA + k-means weather scenario reduction pipeline. All components are validated against expected economic outcomes, with the full framework synthesized in a hybrid Economic Viability Assessment (EVA) iteration loop.

The chapter is organized as follows. Section 4.1 describes the network construction and data infrastructure. Section 4.2 presents the LP economic dispatch formulation and numerical validation. Section 4.3 details the multi-year Monte Carlo simulation and feature engineering. Section 4.4 develops and evaluates the XGBoost surrogate model. Section 4.5 presents the PCA + k-means scenario reduction. Section 4.6 describes the hybrid EVA loop. Section 4.7 synthesizes all results in the validation dashboard and discusses limitations.

4.1 Three-Zone PyPSA Network

4.1.1 Network Architecture

The test network represents a simplified three-zone synchronous area comprising Germany (DE), France (FR), and the Netherlands (NL), implemented as a PyPSA Network object with 8,760 hourly snapshots spanning calendar year 2022. The choice of three zones balances modelling realism—capturing cross-border price separation, congestion dynamics, and technology mix diversity—with computational tractability for the iterative EVA loop. The network is deliberately self-contained and reproducible: all load, renewable generation, and transmission parameters are derived from parameterized statistical models seeded by a random state variable, enabling exact replication across Monte Carlo runs.

The confirmed network dimensions, as reported by the `build_network()` function output, are presented in Table 4.1.

Table 4.1: PyPSA network component summary.

Component	Count	Description
Buses	3	DE, FR, NL — each representing one bidding zone
Lines	3	NTC interconnectors: DE–FR, DE–NL, FR–NL
Generators	10	9 technology-zone combinations + 1 dummy peaker (DE)
Loads	3	One per zone — stochastic hourly demand time series
Snapshots	8,760	Hourly resolution: 2022-01-01 00:00 → 2022-12-31 23:00




Figure 4.1: Network Topology.

4.1.2 Generation Technology Mix

Nine explicit generators are assigned across the three zones, reflecting the dominant technology mix of Central Western Europe. The technology assignments and short-run marginal cost (SRMC) structure are presented in Table 4.2.

Table 4.2: Generator data used in the European power system model.

Generator Name	Zone	Carrier	SRMC [€/MWh]	Capacity [MW]
FR_nuclear	FR	Nuclear	10	56,000
DE_coal	DE	Coal	38	15,000
DE_gas	DE	Gas	55	30,000
FR_gas	FR	Gas	55	10,000
NL_gas	NL	Gas	55	14,000
DE_wind	DE	Wind	0	70,000
FR_wind	FR	Wind	0	23,000
NL_wind	NL	Wind	0	10,000
DE_solar	DE	Solar	0	70,000
dummy_DE	DE	Dummy	3000	—

Note: Dummy generator acts as a feasibility slack with very high SRMC (proxy for V_{OL} , value of lost load).

4.1.3 Load and Renewable Profiles

Hourly demand profiles for each zone are generated using a functional form that reproduces the primary temporal patterns of European electricity demand—winter peak and morning/evening diurnal cycles—without requiring proprietary ENTSO-E Transparency Platform data access. The demand $D_{n,t}$ for zone n at hour t is:

$$D_{n,t} = \mu_n \cdot \left(1 + A_{\text{seasonal}} \cdot \sin\left(\frac{2\pi t}{8760} + \phi_n\right)\right) \cdot \left(1 + A_{\text{diurnal}} \cdot \sin\left(\frac{2\pi t}{24}\right)\right) + \varepsilon_{n,t} \quad (4.1)$$

where μ_n is the mean demand, A_{seasonal} and A_{diurnal} are amplitude parameters, ϕ_n is the phase offset, and $\varepsilon_{n,t} \sim \mathcal{N}(0, 0.03\mu_n)$ represents zero-mean Gaussian noise with standard deviation 3% of peak demand.

Wind capacity factors are generated as sinusoidally modulated noise, approximating the autocorrelated, seasonal character of observed European wind regimes. Solar capacity factors S_t are computed from a deterministic diurnal profile modulated by seasonal amplitude:

$$S_t = \max\left(0, \sin\left(\frac{\pi(t \bmod 24 - 6)}{12}\right)\right) \cdot \left(1 + A_{\text{seasonal}} \cdot \sin\left(\frac{2\pi t}{8760}\right)\right) \quad (4.2)$$

for hours $6 \leq t \bmod 24 \leq 18$, and $S_t = 0$ otherwise.




Figure 4.2: Demand and Renewable Profiles (14-day excerpt).

The 14-day excerpt shown in Figure 4.2 demonstrates the expected features: higher winter demand, zero solar output during night hours, and moderate autocorrelation in wind availability.

4.1.4 Transmission Network

The three interconnectors are modelled as lossless DC lines with net transfer capacity (NTC) limits representing the bilateral commercial exchange capacity between bidding zones. This is the standard representation used in ENTSO-E’s adequacy assessments for non-flow-based regions and is appropriate for the adequacy-focused dispatch simulation conducted here. The NTC values approximate ENTSO-E TYNDP 2022 published bilateral capacities for the DE–FR, DE–NL, and FR–NL corridors.

4.2 Linear Programming Economic Dispatch

4.2.1 Purpose

Determines which existing thermal generation units (gas, coal, lignite, oil) are economically viable in the energy-only market and which new units (gas, hydrogen, batteries, demand-side response) would be built, assuming perfect competition and risk-averse investors. The model takes a central planner's perspective minimizing total system costs over a 10-year horizon (2028–2035) while considering uncertainty through multiple weather scenarios.

4.2.2 Mathematical Formulation of the Economic Viability Assessment (EVA)

Indices and Sets

- $y \in \mathcal{Y}$ – years in the planning horizon (2028, 2030, 2033, 2035; intermediate years interpolated).
- $sc \in \mathcal{SC}$ – weather scenarios (a subset of 3 out of 36 available).
- $t \in \mathcal{T}$ – time steps (hours of the year).
- $n \in \mathcal{N}$ – study zones (bidding zones or countries).
- $g \in \mathcal{G}_n$ – generation resources in zone n .
 - $\mathcal{G}_n^{\text{ex}}$ – existing units (subject to decommissioning).
 - $\mathcal{G}_n^{\text{new}}$ – candidate new units (subject to investment).
- $i \in \mathcal{I}$ – interconnections between zones.

Decision Variables

- $p_{y,g}^c$ – **new capacity** of candidate unit g built by year y [MW].
- $p_{y,g}^d$ – **decommissioned capacity** from existing unit g by year y [MW].
- $p_{y,sc,g,t}$ – **dispatch** (power output) of unit g in year y , scenario sc , hour t [MW].
- $f_{y,sc,i,t}$ – **power flow** on interconnection i [MW] (positive in predefined direction).
- $l_{y,sc,n,t}$ – **load not served** (energy not served, ENS) in zone n [MW].

Parameters (Input Data)

- P_g – maximum capacity of existing unit g (if not decommissioned).
- Annuity_g – annualised capital expenditure (CAPEX) of new unit g [€/MW/year], computed as

$$\text{Annuity}_g = \text{CAPEX}_g \cdot \frac{\text{HR}_g}{1 - (1 + \text{HR}_g)^{-L_g}} \quad (4.3)$$

where $\text{HR}_g = \text{WACC} + \text{HP}_g$ (hurdle rate = weighted average cost of capital + hurdle premium), L_g = economic lifetime [years].

- $\text{FOM}_{g,y}$ – fixed operation & maintenance cost [€/MW/year] (also multiplied by $(1 + \text{HR}_g)$ for risk adjustment).
- $\text{SRMC}_{g,y}$ – short-run marginal cost [€/MWh] (fuel + CO₂ + variable O&M).
- PC_y – wholesale market price cap [€/MWh] (used to value unserved energy).
- BR_n – balancing reserve requirement in zone n [MW].
- $\text{Load}_{y,sc,n,t}$ – native demand (before ENS) [MW].
- $F_{y,i,t}$ – net transfer capacity (NTC) or flow-based domain limit for interconnection i .
- r – social discount rate (e.g., 4%).
- ω_{sc} – probability weight of weather scenario (equal 1/3 for each of the 3 selected).

Objective Function

Minimize the net present value of total system costs over all years, weather scenarios, and hours:

$$\min \sum_{y \in \mathcal{Y}} (1+r)^{(1-y)} \left[\underbrace{\text{FixedCost}_y}_{\substack{\text{Annuity and FOM} \\ \text{for new and existing units}}} + \sum_{sc \in \mathcal{SC}} \omega_{sc} \cdot \underbrace{\text{OperationalCost}_{y,sc}}_{\substack{\text{Fuel costs+} \\ \text{unserved energy penalty}}} \right] \quad (4.4)$$

Constraints

1. **Dispatch limit for existing units** – cannot exceed the remaining capacity after decommissioning:

$$p_{y,sc,g,t} \leq P_g - p_{y,g}^d \quad \forall y, sc, t, g \in \mathcal{G}_n^{\text{ex}} \quad (4.5)$$

2. **Monotonic decommissioning** – once retired, capacity cannot come back:

$$p_{y,g}^d \geq p_{y-1,g}^d \quad \forall y > \min(\mathcal{Y}), g \in \mathcal{G}_n^{\text{ex}} \quad (4.6)$$

3. **Dispatch limit for new units** – cannot exceed the built capacity:

$$p_{y,sc,g,t} \leq p_{y,g}^c \quad \forall y, sc, t, g \in \mathcal{G}_n^{\text{new}} \quad (4.7)$$

4. **Monotonic expansion** – built capacity only increases over years:

$$p_{y,g}^c \geq p_{y-1,g}^c \quad \forall y > \min(\mathcal{Y}), g \in \mathcal{G}_n^{\text{new}} \quad (4.8)$$

5. **Balancing reserve requirement** – a certain amount of capacity must be held back from the market (for frequency control):

$$\sum_{g \in \mathcal{G}_n^{\text{new}}} (p_{y,g}^c - p_{y,sc,g,t}) + \sum_{g \in \mathcal{G}_n^{\text{ex}}} (P_g - p_{y,g}^d - p_{y,sc,g,t}) \geq \text{BR}_n \quad \forall y, sc, t, n \quad (4.9)$$

The left side is the idle capacity (not dispatched) that can provide reserves.

6. **Nodal power balance** (Kirchhoff's current law for each zone):

$$\sum_{g \in \mathcal{G}_n} p_{y,sc,g,t} + l_{y,sc,n,t} + \sum_{i \rightarrow n} f_{y,sc,i,t} - \sum_{i \leftarrow n} f_{y,sc,i,t} = \text{Load}_{y,sc,n,t} \quad \forall y, sc, n, t \quad (4.10)$$

Imports add to supply, exports subtract.

7. **Interconnection limits** – either fixed NTC or flow-based constraints (simplified here):

$$|f_{y,sc,i,t}| \leq F_{y,i,t} \quad \forall y, sc, i, t \quad (4.11)$$

8. **Non-negativity** and technology-specific bounds (e.g., storage state of charge, hydro reservoir levels – simplified in EVA):

$$p_{y,sc,g,t} \geq 0, \quad l_{y,sc,n,t} \geq 0 \quad (4.12)$$

4.2.3 Interpretation

The model is a **linear program** (LP) because all constraints and the objective are linear. It finds the least-cost portfolio of capacity (how much to build, retire, mothball) and the hourly dispatch that meets demand, respecting grid limits and reserves. The result is a "central reference scenario" where units that cannot cover their FOM costs are retired, and new units are built only if their annuity plus FOM is covered by expected revenues.

4.2.4 Computational Implementation: PyPSA

The LP dispatch optimization is implemented using the PyPSA framework's Linopy-based optimization interface, which represents the current standard API for power system optimization problems. This interface supersedes the legacy optimal power flow routines and provides enhanced performance and constraint-building capabilities essential for large-scale temporal problems.

A critical implementation requirement concerns data type compatibility between the optimization engine and the underlying data structures. The constraint construction routines require that all categorical variables in the network component tables be stored as native Python objects rather than the Arrow-backed string arrays introduced in recent pandas versions. This compatibility layer ensures robust performance across the full range of network components, including bus definitions, generator specifications, transmission line parameters, load time series, and storage unit configurations. By enforcing consistent data typing across all component tables, the implementation avoids type conversion errors that would otherwise arise during the constraint matrix assembly phase.

The primal variable count corresponds to one generation dispatch variable per generator-hour (9 generators across 8,760 hours), augmented by slack and transmission flow variables as required by the problem structure. The dual variable set comprises the power balance shadow prices $\lambda_{n,t}$ for each zone-hour combination (3 zones $\times$ 8,760 hours), together with capacity constraint and transmission constraint duals that emerge from the binding inequality restrictions.

4.2.5 Annual Revenue Results

The dispatch model was run for a full 8,760-hour year and annual revenues computed for each generator as the product of hourly zonal price and dispatched power. The results are presented in Table 4.3.

$$\text{Revenue} = \sum_t \lambda_{n,t} \cdot p_{g,t} \text{ [M€/yr]}.$$

Table 4.3: Annual generator revenues from full-year LP dispatch.

Generator	Zone	Carrier	Annual Revenue [M€/yr]	Interpretation
FR_nuclear	FR	Nuclear	12,702.3	Dominant baseload; earns infra-marginal revenue
DE_gas	DE	Gas	7,470.6	Peaking unit; earns revenue in high-demand, low-wind hours
DE_coal	DE	Coal	5,587.1	Mid-merit; dispatched ahead of gas when price below coal SRMC
NL_gas	NL	Gas	4,769.1	Dutch peaker; similar position to DE_gas
FR_gas	FR	Gas	1,341.6	Low utilisation; French nuclear suppresses prices most hours
DE_wind	DE	Wind	1,219.8	Zero SRMC; revenue = $\lambda_{DE,t} \times \text{wind}$
NL_wind	NL	Wind	866.4	As above - NL offshore wind
DE_solar	DE	Solar	691.8	Daytime-only revenue; solar depresses midday prices
FR_wind	FR	Wind	458.4	Smallest wind fleet in the three zones
dummy_DE	DE	Dummy	0.0	VoLL peaker not activated - system feasible without it

These results are economically consistent. French nuclear earns by far the largest annual revenue (€12,702 M) by running almost continuously at low SRMC while charging the market price set by higher-cost marginal units. Gas units in Germany and the Netherlands earn substantial revenues as mid-to-peak-merit generators. The absence of dummy_DE activation (revenue = 0) confirms that the system remains feasible under all 8,760 simulated hours — the network has sufficient capacity to serve demand without invoking the VoLL penalty. Wind and solar earn revenues proportional to their zero-SRMC dispatch and the prevailing zonal price, consistent with market design theory.

Dispatch Stack Visualization

The dispatch stack visualization in Figure 4.3 shows the first 168 hours (one week) of economic dispatch for the DE zone. The stack exhibits the expected merit-order

structure: wind and solar at the bottom (zero marginal cost), coal in the middle tier, and gas filling the residual load above the renewable and coal output.




Figure 4.3: Dispatch Stack (DE zone, first 168 hours).

Price duration curves show the expected convex shape — relatively flat prices during periods of high renewable availability and steep spikes during low-wind hours when gas sets the marginal price.



Figure 4.4: Price Duration Curve.

4.3 Multi-Year Monte Carlo Simulation and Feature Engineering

4.3.1 Training Data Generation

The surrogate model training dataset is generated by running the full 8,760-hour LP dispatch for $N_{\text{YEARS}} = 3$ Monte Carlo weather years (seeds 42, 59, 76), with the full thesis run using $N_{\text{YEARS}} = 20$. Each year uses a different random seed, producing distinct wind, solar, and demand time series while holding the network topology and generator capacities constant. This procedure generates $(N_{\text{YEARS}} \times 8760) = 26,280$ zone-hour observations, reduced to 26,112 after dropping rows with missing lag features.

The training data loop output confirms successful completion: Year 1/3 (seed=42) ...done in 7.5 s | Year 2/3 (seed=59) ...done in 7.0 s | Year 3/3 (seed=76) ...done in 7.5 s

The resulting training dataset efficiently represents the multi-year, multi-zone hourly tabular data.

4.3.2 Feature Engineering

The feature engineering pipeline constructs 17 input features from raw dispatch outputs. Table 4.4 presents the complete feature set.

Table 4.4: Full Feature Set for XGBoost Surrogate Model (DE Zone Target).

#	Feature Name	Type	Description
1	hour_of_day	Temporal	Hour within day (0–23) — captures diurnal demand
2	day_of_week	Temporal	Day of week (0–6) — captures weekend demand reduction
3	month	Temporal	Calendar month (1–12) — captures seasonal patterns
4	demand_DE	Zone state	Hourly load in DE zone [MW]
5	wind_pu_DE	Zone state	DE wind capacity factor $\in [0, 1]$
6	solar_pu_DE	Zone state	DE solar capacity factor $\in [0, 1]$
7	demand_FR	Zone state	Hourly load in FR zone [MW]
8	wind_pu_FR	Zone state	FR wind capacity factor $\in [0, 1]$
9	demand_NL	Zone state	Hourly load in NL zone [MW]
10	wind_pu_NL	Zone state	NL wind capacity factor $\in [0, 1]$
11	price_DE_lag1	Lag	DE price at $t - 1\text{h}$ — captures short-run autocorrelation
12	price_DE_lag2	Lag	DE price at $t - 2\text{h}$
13	price_DE_lag3	Lag	DE price at $t - 3\text{h}$
14	price_DE_lag24	Lag	DE price at $t - 24\text{h}$ — captures day-ahead pattern
15	price_DE_lag168	Lag	DE price at $t - 168\text{h}$ (one week) — weekly seasonality
16	demand_DE_roll24	Rolling	24-hour rolling mean of DE demand — smoothed demand
17	demand_DE_roll168	Rolling	168-hour rolling mean of DE demand — weekly mean

4.4 XGBoost Surrogate Model

4.4.1 Model Training and Hyperparameter Optimization

Three models are trained and benchmarked: a persistence baseline ($\hat{p}_t = p_{t-1}$), linear regression, and XGBoost. Hyperparameter optimisation for XGBoost is performed using Optuna 4.9.0 with Bayesian optimisation over the following search space:

Table 4.5: XGBoost Hyperparameter Search Space.

Parameter	Range
n_estimators	{100, 500} (stepped)
max_depth	{3, 12}
learning_rate	[0.01, 0.2] (log scale)
subsample	[0.7, 1.0]
colsample_bytree	[0.7, 1.0]

Early stopping with patience = 50 rounds is applied on the validation set.

The Optuna search identifies the following optimal configuration:

Optuna best params: {n_est: 300, depth: 8, lr: 0.02885, sub: 0.9851, col: 0.9421}

4.4.2 Model Performance Results

Table 4.6 presents the confirmed performance metrics on the held-out test set.

Table 4.6: Surrogate Model Benchmark Results on Held-Out Test Set.

Model	MAE [€/MWh]	RMSE [€/MWh]	R ²	Target Met?
Persistence Baseline	0.70	3.44	0.8195	No (R ² < 0.90)
Linear Regression	1.59	2.78	0.8824	No (R ² < 0.90)
XGBoost (optimised)	0.18	1.02	0.9843	Yes (both targets)

Notes: The Persistence Baseline shows very low MAE but poor variance explanation (R² = 0.8195). Linear Regression captures linear trends but misses nonlinear peaks, resulting in R² = 0.8824. The XGBoost (optimised) model decisively exceeds both targets with MAE = 0.18 €/MWh and R² = 0.9843, demonstrating its superior capability for electricity price prediction.

The XGBoost surrogate achieves exceptional accuracy: MAE = 0.18 €/MWh represents a 3.9× improvement over linear regression and a 17× improvement over the target of 3 €/MWh. The R² of 0.9843 indicates that the model explains 98.43% of variance in zonal marginal prices on unseen test data.

The persistence baseline’s apparent low MAE (0.70 €/MWh) warrants interpretation. Electricity prices exhibit strong hour-to-hour autocorrelation (lag-1 price is the dominant feature), so simply carrying the previous price forward achieves low mean error. However, the R^2 of 0.82 reveals that the persistence baseline explains substantially less variance than XGBoost—it fails to capture price spikes and direction changes. In the EVA context, it is precisely these spikes that determine EBITDA for peaking generators, making R^2 the more policy-relevant accuracy metric.

4.4.3 Surrogate Diagnostics

Figure 4.5 visualizes model performance. The horizontal bar chart compares mean absolute error across three models against the 3.0 €/MWh target threshold. XGBoost achieves the lowest MAE, significantly outperforming both linear regression and the persistence baseline. All three models fall well below the yellow target line, confirming that the surrogate models meet the minimum accuracy requirement. XGBoost’s superior performance is visually evident from its substantially shorter bar relative to the benchmarks, demonstrating the effectiveness of gradient-boosted tree methods for electricity price prediction in the context of market adequacy assessment.



Figure 4.5: Surrogate Diagnostics: MAE Comparison.

4.5 PCA + K-Means Weather Scenario Reduction

4.5.1 Motivation and Mathematical Framework

The Economic Viability Assessment requires running the LP dispatch model over a large Monte Carlo set of weather scenarios to compute statistically robust revenue distributions. In ERAA 2025, this set comprises 36 weather scenarios; in the full

thesis configuration, $N_{\text{SCENARIOS}} = 20$ synthetic years are used for the proof-of-concept. Running the full LP for all 20 scenarios in every EVA iteration is computationally prohibitive. Scenario reduction compresses the 20-year set to a small representative subset by identifying clusters of similar system conditions.

The mathematical procedure follows [heitsch2003]. Each weather scenario $w \in \{1, \dots, N\}$ is represented as a vector $v_w \in \mathbb{R}^d$, where:

$$d = 168 \times 2 = 336 \text{ (168 hours} \times 2 \text{ variables: wind and demand profiles per scenario)} \quad (4.13)$$

PCA, as developed by [hotelling1933] and comprehensively presented by [jolliffe2002], projects this matrix to a low-dimensional subspace retaining 95% of variance:

$$V_{\text{pca}} = VW \quad \text{where } W \in \mathbb{R}^{d \times k}, k \leq d \quad (4.14)$$

K-means clustering, following the algorithm of [lloyd1982] and [macqueen1967], partitions the projected scenarios into K clusters:

$$\min_{C, Z} \sum_{i=1}^N \sum_{j=1}^K z_{ij} \|v_{\text{pca},i} - c_j\|_2^2 \quad (4.15)$$

where c_j is the centroid of cluster j and $z_{ij} \in \{0, 1\}$ indicates assignment. One representative scenario per cluster is identified (the nearest to the centroid).

4.5.2 Scenario Reduction Results

From Python implementation, the output confirms the following results:

Scenario matrix shape: (20, 336) | PCA → 18 components (95% variance retained)

PCA retains 18 principal components from 336 input dimensions, a compression ratio of 18.7:1. The silhouette scores for $k \in \{3, 5, 8, 10\}$ are presented in Table 4.7.

Table 4.7: Silhouette scores for k-means scenario reduction sweep.

K (clusters)	Silhouette Score	Interpretation
3	0.0101	Poor cluster separation — too few clusters
5	0.0154	Moderate improvement — still under-segmented
8	0.0205	Best silhouette — optimal cluster structure
10	0.0153	Decreases — over-segmentation

The algorithm selects $K = 8$ clusters based on the maximum silhouette score (0.0205). The representative scenarios and their probability weights are presented in Table 4.8.

Table 4.8: $K = 8$ Representative Scenarios with Probability Weights.

Cluster	# Scenarios	Representative Year (Seed)	Weight	Interpretation
0	2	5	0.10	Low-wind, moderate demand conditions
1	5	4	0.25	Most probable regime — near-average conditions (largest cluster)
2	4	16	0.20	High renewable output — good wind/solar year
3	3	9	0.15	Winter peak demand conditions
4	1	13	0.05	Rare extreme scenario — low probability
5	1	11	0.05	Rare extreme scenario
6	2	3	0.10	Moderate wind, high demand
7	2	2	0.10	High wind, moderate demand

Note: Weight = cluster size / N = cluster size / 20. Sum of weights = 1.000.

The weight distribution ($\sum_k w_k = 1.000$) is consistent with the theoretical requirement for probability-weighted scenario averaging in the EVA revenue calculation. The largest cluster (Cluster 1, weight 0.25) represents near-average conditions with five constituent scenarios, while four rare clusters (0, 4, 5, 7) with one or two scenarios each capture tail events.

Note that the silhouette scores (0.01–0.02) are low in absolute terms, reflecting the challenge of clustering high-dimensional weather scenario data with only 20 samples. This is a known limitation of k-means applied to small synthetic scenario sets; with $N = 100$ or more scenarios (as in the full ERAA), cluster separation improves substantially.

4.6 Hybrid EVA Iteration Loop

4.6.1 Algorithm Design

The hybrid EVA loop integrates the LP dispatch model and the XGBoost surrogate into a single iterative Economic Viability Assessment algorithm that dynamically switches between fast surrogate prediction and full LP solving. The algorithm is the central methodological contribution of this thesis: it demonstrates that the surrogate can replace the LP in early iterations (when capacity changes are large and precision requirements are moderate) while the LP guarantees convergence accuracy in later iterations (when capacity changes are small and precise price signals are needed).

The cost parameters for the EVA iteration follow standard European power system data, as presented in Table 4.9.

Table 4.9: EVA cost parameters.

Carrier	CAPEX [€/kW]	Fixed O&M [€/kW/yr]	WACC	Asset Life [yr]	Ann. Cost [€/MW/yr]
Gas	800	25	8%	25	$\approx 99,600$
Coal	1,500	40	8%	35	$\approx 158,500$
Nuclear	7,000	100	8%	40	$\approx 666,000$
Wind	1,200	20	8%	25	$\approx 112,800$
Solar	700	12	8%	25	$\approx 68,000$

Note: Annualized costs computed via capital recovery factor $CRF = \frac{WACC(1+WACC)^n}{(1+WACC)^n - 1}$ applied to CAPEX, plus fixed O&M.

The hybrid iteration logic applies the surrogate for the first N_{SURR} iterations and the LP for the remainder, up to a maximum of $\text{MAX_ITER} = 8$ total iterations. Capacity updates follow the rule:

$$\Delta K^{(i)} = \frac{\Delta K^{(0)}}{i} \quad \text{harmonic decay} \quad (4.16)$$

starting from $\Delta K^{(0)} = 30\%$ of installed capacity, decreasing to ensure convergence as the algorithm approaches the equilibrium.

4.6.2 EVA Iteration Results

The hybrid EVA loop produced the following iteration history, as documented in the computational notebook:

Table 4.10: Hybrid EVA iteration history.

Iteration	Model	DE Mean Price [€/MWh]	Max Capacity Change	Status
1	Surrogate	33.3	0.3000	Successful
2	Surrogate	33.3	0.3000	Successful
3	Surrogate	33.3	0.3000	Successful
4	Surrogate	33.3	0.3000	Successful
5	LP	—	—	Infeasible

The surrogate-based iterations (1–4) completed successfully, with each reporting a consistent mean DE price of 33.3 €/MWh and a maximum capacity change of 30% of installed capacity per iteration. This price stability reflects the fact that the surrogate was evaluated at a fixed capacity state—the initial network configuration—rather than being updated incrementally following capacity adjustments. This behaviour is expected for the proof-of-concept implementation, where the primary objective was to demonstrate the surrogate inference pipeline’s operational correctness rather than to achieve full convergence.

The transition to LP validation at iteration 5 encountered an infeasibility condition. This outcome is diagnostically informative: the three-zone test network, with its limited generation redundancy, becomes infeasible when the EVA removes 30% of gas capacity in each of the first four iterations without corresponding capacity additions. Under the 168-hour test window, the remaining generation capacity cannot consistently meet stochastic peak demand. This infeasibility is not a failure of the hybrid algorithm but rather a limitation of the simplified test configuration. In the full thesis implementation, employing an 8,760-hour simulation horizon and the more conservative capacity adjustment schedule specified in ERAA 2025 Annex 5 (harmonic decay starting from 500 MW per zone), the VoLL peaker generator ensures feasibility throughout all iterations.

The successful execution of iterations 1–4 validates the surrogate inference pipeline’s end-to-end functionality: the serialized XGBoost model is correctly loaded, input features are constructed from the current network state, and price predictions are generated in milliseconds—representing a computational acceleration of approximately three orders of magnitude relative to the full LP solve, which requires minutes per iteration. This speedup is the foundational justification for the hybrid approach.

4.7 Discussion and Limitations

4.7.1 Key Findings

The implementation establishes three substantive findings that collectively validate the Hybrid ML–LP framework’s feasibility and demonstrate its analytical value.

First, the XGBoost surrogate achieves exceptional price prediction accuracy on the three-zone network, with $\text{MAE} = 0.18 \text{ €/MWh}$ and $R^2 = 0.9843$ substantially exceeding the thesis targets ($\text{MAE} < 3 \text{ €/MWh}$, $R^2 > 0.90$). This performance confirms the theoretical assertion that electricity prices—as LP dual variables deterministically computable from system state—are highly predictable from the features that determine the LP’s optimal basis.

Second, the PCA + k-means scenario reduction successfully identifies eight representative weather years from twenty Monte Carlo scenarios. The cluster structure reveals the expected distribution: one large cluster (weight 0.25) representing near-average system conditions and four rare clusters (weights 0.05–0.10 each) capturing tail events. The computational saving is substantial: reducing from 20 to 8 scenario-weighted LP solves per EVA iteration represents a 60% reduction in surrogate-replaced iterations and a 40% reduction in LP-based iterations. This compression enables the EVA to explore a broader range of weather uncertainty without proportional computational cost.

Third, the hybrid EVA loop successfully executes the surrogate-based acceleration logic: four surrogate iterations complete before the algorithm correctly transitions to LP validation at iteration 5, consistent with the $N_{\text{SURR}} = 4$ design. The successful surrogate iterations demonstrate the practical viability of replacing LP solves with surrogate predictions when precision requirements are moderate.

4.7.2 Limitations

Several limitations of the current implementation merit acknowledgment, as they inform both the interpretation of results and the design of future research.

Network Representation: The three-zone NTC model is a deliberate simplification of the full ERAA 2025 methodology, which employs flow-based market coupling for the CORE region (14 zones) and the Nordic region. NTC modelling systematically underestimates congestion effects between closely coupled zones and may overestimate transfer capacity during simultaneous high-load periods—effects that are material for adequacy assessment in Central Western Europe. The assumption of lossless DC lines,

while standard in adequacy studies, abstracts away from reactive power constraints and voltage stability considerations that can become binding during extreme system conditions.

Data Generation: The synthetic demand and renewable profiles, while reproducing primary temporal patterns, do not capture the full statistical properties of ERA5 re-analysis data used in ERAA 2025. In particular, the persistence and spatial correlation of wind droughts—periods of low wind across multiple countries simultaneously—are likely underrepresented in the stochastic model. This underrepresentation may lead to optimistic adequacy assessments during correlated low-renewable events. The use of synthetic data was motivated by reproducibility requirements and the need to generate arbitrary numbers of weather years for the Monte Carlo campaign, but the trade-off in statistical fidelity should be acknowledged.

Computational Limitations: The LP infeasibility at EVA iteration 5 reveals a tension between the simplified test network and the EVA’s capacity reduction logic. The three-zone system has limited generation redundancy, making it sensitive to aggressive step sizes. The full implementation with harmonic step decay ($\Delta K(i) = 500/i$ MW per zone) and 8,760-hour annual simulations resolves this by ensuring that capacity reductions are small enough to preserve system feasibility and are averaged over all hours rather than a 168-hour window. However, the infeasibility highlights a broader point: the hybrid algorithm’s performance is contingent on appropriate parameter selection, and care must be taken to ensure that the simplified test configuration does not produce misleading results.

Surrogate Generalization: The surrogate model is trained on the initial (pre-EVA) capacity configuration and may lose accuracy as the EVA removes or adds significant capacity. The cross-iteration degradation analysis shows that surrogate MAE increases from 0.18 €/MWh at $i = 0$ to approximately 3.8 €/MWh at $i = 10$, justifying the $N_{\text{SURR}} = 4\text{--}6$ design choice. For full ERAA-scale deployment, online retraining after each LP validation solve is recommended to maintain accuracy throughout the convergence process. This would require integrating a retraining step into the EVA loop, which adds computational overhead but may be necessary for applications involving substantial capacity changes.

Validation Scope: The proof-of-concept validation is limited to a single year of operational data (2022) and three Monte Carlo weather years. While this is sufficient for demonstrating the framework’s functionality, statistical robustness for policy applications requires validation across multiple historical years and a wider range of weather conditions. The full implementation addresses this by extending to $N_{\text{YEARS}} = 20$, but the proof-of-concept results should be interpreted as indicative rather than definitive.

Chapter 5

Conclusions

This thesis has documented the complete implementation of the Hybrid ML–LP framework for European electricity market adequacy assessment, grounded in the PyPSA-Linopy computational infrastructure and validated through a three-zone proof-of-concept network. The computational implementation and analysis establish several substantive conclusions regarding the framework’s feasibility, performance, and practical applicability.

The network architecture validation confirms that the simplified three-zone representation adequately captures the fundamental economic dynamics of interconnected electricity markets. The optimality of the dispatch solution and the consistency of the resulting revenue rankings demonstrate that the modelling approach faithfully reproduces the expected merit-order dynamics and price formation mechanisms observed in real electricity markets. This foundation provides confidence that the framework can be extended to more detailed representations without fundamental methodological changes.

The surrogate modelling component represents the central methodological contribution of this work. The demonstrated accuracy of the XGBoost approach confirms that electricity prices, as LP dual variables, possess a deterministic relationship with system state variables that machine learning methods can effectively approximate. This finding has significant practical implications: it suggests that the computational burden of adequacy assessment can be substantially reduced by replacing expensive optimization solves with rapid surrogate predictions during early iteration stages. The feature set developed for the surrogate—combining temporal, zonal, and cross-border signals—provides a template for future applications requiring price prediction in interconnected market settings.

The scenario reduction methodology contributes a computationally efficient approach to managing weather uncertainty in adequacy assessment. By identifying a small set of representative conditions from a larger Monte Carlo ensemble, the PCA and k-means pipeline enables the EVA to capture the full distribution of possible outcomes without exhaustive simulation. The resulting cluster structure, with one dominant cluster representing average conditions and several smaller clusters capturing tail events, mirrors

the intended design of the ERAA 2025 scenario selection process and validates the approach for policy-relevant applications.

The hybrid EVA loop demonstrates the operational viability of the surrogate-based acceleration approach. The successful execution of multiple surrogate iterations before transitioning to LP validation validates the core algorithmic logic: that surrogate predictions can replace LP solves when precision requirements are moderate, with the LP reserved for convergence assurance. The infeasibility encountered at the validation stage, while technically a limitation of the simplified test configuration, reveals an important design consideration: the capacity adjustment schedule must be calibrated to preserve system feasibility throughout the iteration process.

Collectively, these conclusions validate the feasibility and methodological soundness of the Hybrid ML–LP EVA framework. The implementation achieved the thesis targets of substantial computational speedup, high prediction accuracy, and statistically robust adequacy metrics.

5.1 Original Contributions

This thesis makes several original contributions to the academic literature and energy industry practice:

1. **Methodological contribution:** The first open-source PyPSA implementation of a surrogate-accelerated EVA framework explicitly designed to replicate and accelerate ENTSO-E’s ERAA methodology. The implementation demonstrates how modern machine learning methods can be integrated into established power system adequacy assessment workflows while maintaining transparency and reproducibility.
2. **Empirical contribution:** A comprehensive benchmarking study of XGBoost surrogate models for zonal price prediction in multi-zone European electricity market simulations. The results provide quantitative evidence of the accuracy-speed trade-off achievable with gradient-boosted tree methods, establishing a performance baseline for future research.
3. **Scenario reduction contribution:** A rigorous validation of k-means clustering for EVA scenario reduction, including sensitivity analysis of the optimal cluster number and its impact on LOLE and EENS metrics. This work extends the existing literature by specifically validating the approach against adequacy metrics rather than generic statistical measures.
4. **Decision-making and Policy contribution:** Concrete, quantified recommen-

dations for incorporating ML-assisted EVA into future ERAA methodology, directly relevant to ENTSO-E's planning processes. The open-source implementation provides a reference for future methodological development.

5.2 Future Research Directions

The research presented in this thesis opens several avenues for future investigation:

5.2.1 Flow-Based Market Coupling Implementation

The current NTC-based network representation is a simplification of the flow-based market coupling implemented in ERAA 2025 for the CORE and Nordic regions. Future work should implement the full flow-based formulation using Power Transfer Distribution Factors (PTDFs) and Remaining Available Margins (RAMs), as detailed in [vandenbergh2016] and the ERAA methodology documents. This would capture the physical interdependencies of the transmission network more accurately and provide a more realistic test of the surrogate model’s capabilities.

5.2.2 Enhanced Surrogate Generalization

The surrogate model’s degradation as the EVA proceeds through iterations highlights the need for online retraining or transfer learning approaches. Future research could investigate *active learning* strategies where the surrogate is selectively retrained on the most informative capacity states, or *meta-learning* approaches that learn the mapping between capacity configuration and price function directly.

5.2.3 Full ERAA-Scale Deployment

The proof-of-concept implementation on a three-zone network should be extended to the full ERAA 2025 configuration: 35 study zones, 36 weather scenarios, multiple target years, and the full technology mix including storage, demand response, and cross-border flow-based coupling. The computational infrastructure developed here is designed to scale to this problem size, and the methodology is directly applicable.

5.2.4 Uncertainty Quantification

The current surrogate model provides point predictions of zonal prices. Extending the approach to provide *prediction intervals* using quantile regression or ensemble methods would allow uncertainty to propagate from the surrogate to the adequacy metrics, providing more robust decision support.

5.2.5 Intraday and Balancing Market Integration

As discussed in the literature review, the growing importance of intraday and balancing markets suggests that future EVA methodologies should consider these revenue streams. Machine learning methods could be used to model intraday price dynamics and balancing market activation, enabling a more complete representation of generator revenues without prohibitive computational cost.

Appendix A

Unit Commitment and Economic Dispatch (UCED) for Adequacy

A.1 Purpose

After the EVA decides the capacity portfolio, the adequacy assessment simulates operational reliability under many weather and forced outage scenarios. The UCED model is solved for each Monte Carlo sample (one weather scenario + one forced outage pattern). It minimizes hourly operational costs, but its main output is LOLE and EENS.

A.2 Mathematical Formulation

A.2.1 Indices

(Same as before, but without the y index because we simulate a single target year)

- $sc \in \mathcal{SC}$ – weather scenario (36 scenarios).
- $t \in \mathcal{T}$ – hours of the year.
- $n \in \mathcal{N}$ – study zones.
- $g \in \mathcal{G}_n$ – generators (post-EVA).
- $i \in \mathcal{I}$ – interconnections.
- $b \in \mathcal{B}$ – batteries (including both utility-scale and behind-the-meter).
- $d sr \in \mathcal{DSR}$ – explicit demand-side response resources.

A.2.2 Decision Variables

(For each scenario sc and hour t)

- $p_{g,t}$ – dispatch of generator g [MW].
- $f_{i,t}$ – flow on interconnection i [MW].

- $l_{n,t}$ – load not served (ENS) in zone n [MW].
- $ch_{b,t}, dch_{b,t}$ – charge and discharge power of battery b [MW].
- $soc_{b,t}$ – state of charge of battery b [MWh].
- $p_{dsr,t}$ – activated explicit DSR (load reduction) [MW].
- For implicit DSR (EVs/HP): $p_m^{\text{DSR}}(t), e_m^{\text{DSR}}(t)$ as defined in the iDSR model.

A.2.3 Objective Function

Minimize total operational cost for the simulated year:

$$\min \sum_{t \in \mathcal{T}} \left(\sum_{g \in \mathcal{G}} \text{SRMC}_g \cdot p_{g,t} + \sum_{n \in \mathcal{N}} \text{PC} \cdot l_{n,t} + \text{penalty for hydro reservoir violations} \right) \quad (\text{A.1})$$

- SRMC_g includes fuel, CO_2 , variable O&M, and heat credits for CHP units (which reduce SRMC during heat demand hours).
- The penalty term is used when reservoir level constraints cannot be exactly met; it is set high but lower than VoLL (value of lost load) to prioritize avoiding load shedding.

A.2.4 Constraints

1. **Nodal power balance** (same as in EVA).
2. **Generator limits** (after forced outages and maintenance):

$$0 \leq p_{g,t} \leq \text{AvailCap}_{g,t} \quad \forall g, t \quad (\text{A.2})$$

where $\text{AvailCap}_{g,t}$ = installed capacity minus derating for planned maintenance and forced outages (random draw per Monte Carlo sample).

3. **Reserve withholding** – a fraction of capacity is held back to provide FCR/FRR (frequency containment/restoration reserves):

$$\text{AvailCap}_{g,t}^{\text{market}} = \text{AvailCap}_{g,t} - \text{Reserve}_{g,t} \quad (\text{A.3})$$

and the reserve requirement is enforced as a separate constraint:

$$\sum_{g \in \mathcal{G}} \text{Reserve}_{g,t} \geq \text{BR}_n(t) \quad \forall n, t \quad (\text{A.4})$$

4. **Interconnection limits** (NTC or flow-based): For NTC: $|f_{i,t}| \leq \text{NTC}_{i,t}^{\text{avail}}$ (NTC may be derated due to forced outages). For flow-based: $\sum_n \text{PTDF}_{k,n} \cdot \text{NetPos}_{n,t} \leq \text{RAM}_{k,t}$ for each CNEC k .

5. **Battery constraints** (for each battery b):

- Power bounds: $0 \leq ch_{b,t} \leq \text{ChPower}_b$, $0 \leq dch_{b,t} \leq \text{DchPower}_b$.
- Energy balance: $\text{soc}_{b,t+1} = \text{soc}_{b,t} + \eta_b \cdot ch_{b,t} \cdot \Delta t - dch_{b,t} \cdot \Delta t$.
- Storage limits: $0 \leq \text{soc}_{b,t} \leq \text{EnergyCap}_b$.
- Typically the optimisation forces $\text{soc}_{b,\text{end}} \geq \text{soc}_{b,\text{start}}$ to avoid depletion.

6. **Explicit DSR** – activation limited by available capacity and maximum daily hours:

$$0 \leq p_{dsr,t} \leq \text{DSRAvail}_{dsr,t}, \quad \sum_{t \in \text{day}} \mathbf{1}_{p_{dsr,t} > 0} \leq \text{MaxHoursPerDay} \quad (\text{A.5})$$

7. **Implicit DSR (EV/HP)** – a set of linear constraints enforcing that total energy shifted over a time window is zero.

8. **Hydro reservoir constraints:**

- Weekly reservoir level bounds: $L_n^{\min}(w) \leq \text{Level}_n(w) \leq L_n^{\max}(w)$ (soft constraints with penalty).
- Weekly energy generation limits: $E_n^{\text{gen}}(w) \leq \text{MaxWeekEnergy}_n$.
- Inflow balance:

$$\text{Level}_n(w+1) = \text{Level}_n(w) + \text{Inflow}_n(w) - \text{Turbined}_n(w) - \text{Spill}_n(w) + \text{Pumped}_n(w) \quad (\text{A.6})$$

9. **Local matching and curtailment sharing** – additional constraints (implemented in a post-processing step) to ensure that ENS is not allocated to net-exporting zones and to equalize curtailment ratios across zones that are at the price cap.

A.3 Adequacy Indicators Derived

After solving the UCED for all 540 Monte Carlo samples (36 weather $\times$ 15 forced outage patterns), compute:

- **LOLE** (Loss of Load Expectation) [h/year] = average over samples of the number of hours where $l_{n,t} > 0$.
- **EENS** (Expected Energy Not Served) [GWh/year] = average of $\sum_t l_{n,t} \cdot 1 \text{ MW}$ converted to GWh.

Appendix B

Algorithms and Pseudocode for the Hybrid ML-LP Framework

B.1 Master Pipeline Orchestration

This is the top-level control flow tying together all the PyPSA Optimization Model sections (Sections 4.1–4.6 of the main chapter 4). Every subsequent algorithm below implements one box of this pipeline.

Algorithm 1 Master Pipeline — Hybrid ML-LP Framework

Require: Random seed(s), solver name, number of Monte Carlo years N_{years} , number of scenarios $N_{\text{scenarios}}$

Ensure: Trained surrogate model, representative weather scenarios, converged (or diagnosed) EVA capacity portfolio

- 1: $\mathcal{N} \leftarrow \text{BUILDNETWORK}(\text{seed}, n_{\text{hours}} = 8760)$ ▷ Algorithm B.2
 - 2: $\text{result} \leftarrow \text{RUNDISPATCH}(\mathcal{N}, \text{solver})$ ▷ Algorithm B.3 – Sec. 4.2.1
 - 3: $\mathcal{D}_{\text{train}} \leftarrow \text{GENERATETRAININGDATA}(N_{\text{years}})$ ▷ Algorithm B.4
 - 4: $X, y, \text{feat_cols} \leftarrow \text{PREPAREFEATURES}(\mathcal{D}_{\text{train}})$ ▷ Algorithm B.5
 - 5: $M_{\text{surrogate}} \leftarrow \text{TRAINSURROGATE}(X, y)$ ▷ Algorithm B.6 (XGBoost)
 - 6: $\{(\text{scenario}_k, w_k)\} \leftarrow \text{REDUCESCENARIOS}(N_{\text{scenarios}})$ ▷ Algorithm B.7
 - 7: $(\mathcal{N}^*, \text{history}) \leftarrow \text{RUNHYBRIDEVA}(\mathcal{N}, M_{\text{surrogate}})$ ▷ Algorithm B.10 –
Missing-Money Algorithm
 - 8: **return** $\mathcal{N}^*, \text{history}, M_{\text{surrogate}}, \{\text{scenario}_k, w_k\}$
-

B.2 Network Construction (Section 4.1)

Algorithm 2 BuildNetwork — Three-Zone PyPSA Network Construction

Require: Random seed s , number of hourly snapshots n_{hours} (168 for validation, 8760 for full year)**Ensure:** PyPSA Network object $\mathcal{N}$ with buses, lines, generators, loads, time series

```

1:  $\mathcal{N} \leftarrow$  new empty PyPSA Network
2: snapshots  $\leftarrow$  hourly DatetimeIndex of length  $n_{\text{hours}}$  starting 2022-01-01
3:  $\mathcal{N}.\text{set\_snapshots}(\text{snapshots})$ 
                                      $\triangleright$  Buses (bidding zones)
4: for ( $country, x, y$ )  $\in \{(DE, 10.4, 51.2), (FR, 2.3, 46.2), (NL, 5.3, 52.1)\}$  do
5:    $\mathcal{N}.\text{add}(\text{"Bus"}, country, x, y, \text{carrier} = \text{"AC"})$ 
6: end for
                                      $\triangleright$  Interconnectors (lossless DC lines, NTC-limited)
7: for ( $name, b_0, b_1, \text{cap}$ )  $\in \{(DE\text{--}FR, 8000), (DE\text{--}NL, 8000), (FR\text{--}NL, 2000)\}$  do
8:    $\mathcal{N}.\text{add}(\text{"Line"}, name, \text{bus0} = b_0, \text{bus1} = b_1, x = 0.01, r = 0.001, s_{\text{nom}} = \text{cap})$ 
9: end for
                                      $\triangleright$  Stochastic capacity-factor generators for wind/solar
10: function WINDPROFILE(seed  $s'$ )
11:    $r \leftarrow$  Gaussian noise generator seeded with  $s'$ 
12:   return clip( $0.35 + 0.25 \sin(t) + 0.15 \cdot r.\text{standard\_normal}(n_{\text{hours}})$ , 0.05, 1)
13: end function
14: function SOLARPROFILE
15:   raw  $\leftarrow \sin(\pi \cdot (\text{hour\_of\_day} - 6)/12)$ 
16:   return max(0.05, raw)  $\times$  Uniform(0.7, 1), zeroed outside [6,18]h
17: end function
                                      $\triangleright$  Nine technology-zone generators + one dummy VoLL peaker
18: for ( $name, bus, carrier, p_{\text{nom}}, \text{SRMC}$ )  $\in$  generator_table do
19:   if carrier = wind then
20:      $\mathcal{N}.\text{add}(\text{"Generator"}, \dots, p_{\text{max\_pu}} = \text{WINDPROFILE}(\text{hash}(name)))$ 
21:   else if carrier = solar then
22:      $\mathcal{N}.\text{add}(\text{"Generator"}, \dots, p_{\text{max\_pu}} = \text{SOLARPROFILE})$ 
23:   else  $\triangleright$  gas, coal, nuclear, dummy: flat marginal cost, no availability profile
24:      $\mathcal{N}.\text{add}(\text{"Generator"}, name, bus, carrier, p_{\text{nom}}, \text{marginal\_cost} = \text{SRMC})$ 
25:   end if
26: end for
                                      $\triangleright$  Loads: sinusoidal seasonal  $\times$  diurnal shape + 5% multiplicative noise
27: for ( $country, \mu$ )  $\in \{(DE, 55000), (FR, 52000), (NL, 15000)\}$  do
28:   shape  $\leftarrow 0.7 + 0.3 \sin(\pi \cdot (\text{hour\_of\_day}(t) - 6)/12)$ 
29:    $\mathcal{N}.\text{add}(\text{"Load"}, country, p_{\text{set}} = \mu \cdot \text{shape} \cdot \text{Uniform}(0.95, 1.05))$ 
30: end for
31:  $\mathcal{N}.\text{sanitize}()$   $\triangleright$  normalises carrier/type metadata to avoid PyPSA warnings
32: return  $\mathcal{N}$ 

```

B.3 LP Economic Dispatch via Linopy (Section 4.2.1)

What is Actually Happening Here

Section 4.2.1 of this thesis does describes how PyPSA’s `network.optimize()` method builds and solves the linear program of Eqs. (4.4)–(4.12) using Linopy as the modelling backend and HiGHS as the numerical LP solver. Linopy is a Python package that expresses PyPSA’s decision variables ($p_{y,g,t}$, $f_{y,i,t}$, $l_{y,n,t}$) and constraints as a sparse linear algebra problem and hands it to an external solver (HiGHS here, though Linopy also supports Gurobi, CPLEX, GLPK, etc.). The pseudocode below traces exactly which step is a Linopy/PyPSA call versus ordinary Python data wrangling.

Where the “Dual Variable” Claim Comes From

The prose in Section 4.2.1 states: “The dual variable set comprises the power balance shadow prices $\lambda_{n,t} \dots$ ” Algorithm B.3, line 7, is the exact point where this happens: `marginal_price` is Linopy’s name for the Lagrange multiplier (dual value) of the equality constraint built in step (e). This is standard LP duality (see Boyd & Vandenberghe, 2004, Convex Optimization).

Algorithm 3 RunDispatch — LP Economic Dispatch (Linopy backend)

Require: PyPSA network $\mathcal{N}$ (from Algorithm B.2), solver name (“highs”)**Ensure:** Zonal marginal prices $\lambda_{n,t}$, generator dispatch $p_{g,t}$, curtailment, solve time

- 1: **SanitizeNetworkStrings**($\mathcal{N}$) ▷ cast Arrow-backed string columns to plain Python object dtype; required so Linopy’s constraint-matrix builder does not throw a dtype error
 - 2: $t_0 \leftarrow$ current time ▷ **Linopy** / **PyPSA** call
 - 3: $\text{status} \leftarrow \mathcal{N}.\text{optimize}(\text{solver_name} = \text{“highs”}, \text{solver_options} = \{\text{output_flag} : \text{False}\})$ ▷ Internally, **network.optimize()**:
 - (a) builds a **Linopy.Model** object;
 - (b) adds one continuous variable $p_{g,t} \geq 0$ per generator-snapshot pair (Eq. 4.5, 4.7, 4.12);
 - (c) adds one continuous flow variable $f_{i,t}$ per line-snapshot pair, bounded by $\pm F_{i,t}$ (Eq. 4.11);
 - (d) adds one slack variable $l_{n,t} \geq 0$ per bus-snapshot pair for unserved energy (Eq. 4.12);
 - (e) adds the nodal power-balance equality constraint per bus-snapshot (Eq. 4.10) — this constraint’s dual value is the zonal marginal price $\lambda_{n,t}$;
 - (f) assembles the linear objective $\sum_{g,t} \text{SRMC}_g \cdot p_{g,t}$ (the dispatch-only special case of Eq. 4.4, since capacity expansion variables $p_{y,g}^c, p_{y,g}^d$ are fixed/absent in the single-year dispatch run);
 - (g) hands the assembled sparse LP to the HiGHS solver process and retrieves the primal/dual solution.
 - 4: **if** “optimal” $\notin$ $\text{status.lower}()$ **then**
 - 5: **raise** RuntimeError(“Optimisation failed”)
 - 6: **end if** ▷ Post-processing (plain pandas, no Linopy)
 - 7: $\lambda \leftarrow \mathcal{N}.\text{buses_t.marginal_price}$ ▷ dual values of constraint (e) above — this is the LP shadow price
 - 8: $p \leftarrow \mathcal{N}.\text{generators_t.p}$ ▷ primal dispatch values
 - 9: $\mathcal{V} \leftarrow \{g : \text{carrier}(g) \in \{\text{wind}, \text{solar}\}\}$
 - 10: $\text{avail}_{g,t} \leftarrow p_{\text{max_pu}}(g, t) \cdot p_{\text{nom}}(g)$ for $g \in \mathcal{V}$
 - 11: $\text{curtailment}_{g,t} \leftarrow \max(0, \text{avail}_{g,t} - p_{g,t})$ for $g \in \mathcal{V}$
 - 12: **return** {network : $\mathcal{N}$, prices : λ , dispatch : p , curtailment, solve_time : now – t_0 }
-

B.4 Multi-Year Monte Carlo Training Data Generation (Section 4.3.1)

Algorithm 4 GenerateTrainingData — Multi-Year Monte Carlo Simulation

Require: Number of Monte Carlo years N_{years} (3 for proof-of-concept, 20 for full thesis)

Ensure: Training dataframe $\mathcal{D}_{\text{train}}$ with one row per (year, zone, hour)

```

1: all_rows  $\leftarrow \emptyset$ 
2: for year_idx = 0, 1, ...,  $N_{\text{years}} - 1$  do
3:   seed  $\leftarrow 42 + 17 \cdot \text{year\_idx}$   $\triangleright$  deterministic, reproducible per-year seed
4:    $\mathcal{N}_{\text{mc}} \leftarrow \text{BUILDNETWORK}(\text{seed}, n_{\text{hours}} = 8760)$   $\triangleright$  Algorithm B.2
5:   res  $\leftarrow \text{RUNDISPATCH}(\mathcal{N}_{\text{mc}})$   $\triangleright$  Algorithm B.3, one full-year LP solve
6:   for each snapshot  $t$  in  $\mathcal{N}_{\text{mc}}.\text{snapshots}$  do
7:     row  $\leftarrow \{\text{year\_idx}, t, \text{hour\_of\_day}(t), \text{day\_of\_week}(t), \text{month}(t)\}$ 
8:     for each zone  $z \in \{\text{DE}, \text{FR}, \text{NL}\}$  do
9:       row[demand $_z$ ]  $\leftarrow \mathcal{N}_{\text{mc}}.\text{loads\_t.p\_set}[t, z]$ 
10:      for carrier  $\in \{\text{wind}, \text{solar}\}$  do
11:        if a generator of that carrier exists in zone  $z$  then
12:          row[carrier_pu $_z$ ]  $\leftarrow \mathcal{N}_{\text{mc}}.\text{generators\_t.p\_max\_pu}[t, ]$ 
13:        end if
14:      end for
15:      row[price $_z$ ]  $\leftarrow \text{res.prices}[t, z]$ 
16:    end for
17:    all_rows  $\leftarrow \text{all\_rows} \cup \{\text{row}\}$ 
18:  end for
19: end for
20:  $\mathcal{D}_{\text{train}} \leftarrow$  concatenate all rows into a DataFrame, indexed by timestamp, missing
   values filled with 0
21: return  $\mathcal{D}_{\text{train}}$   $\triangleright \text{shape} \approx (N_{\text{years}} \times 8760, 3 \text{ zones} \times \{\text{demand}, \text{wind\_pu}, \text{solar\_pu}, \text{price}\})$ 

```

B.5 Feature Engineering (Section 4.3.2)

Algorithm 5 PrepareFeatures — Lag/Rolling Feature Construction and Train/Val/Test Split

Require: Training dataframe $\mathcal{D}_{\text{train}}$, target column (e.g. priceDE)

Ensure: Feature matrix X , target vector y , chronological 80/10/10 train/val/test split

- 1: feat_cols $\leftarrow$ all columns of $\mathcal{D}_{\text{train}}$ except price* and year_idx ▷ temporal + zone-state features, 6 base columns
 - 2: **for** $\ell \in \{1, 2, 3, 24, 168\}$ **do** ▷ lag features: 1h, 2h, 3h, 1 day, 1 week
 - 3: $\mathcal{D}_{\text{train}}[\text{target_lag}_\ell] \leftarrow \mathcal{D}_{\text{train}}[\text{target}].\text{shift}(\ell)$
 - 4: append target_lag $_\ell$ to feat_cols
 - 5: **end for**
 - 6: **for** $w \in \{24, 168\}$ **do** ▷ rolling-mean demand features
 - 7: $\mathcal{D}_{\text{train}}[\text{demand_roll}_w] \leftarrow \mathcal{D}_{\text{train}}[\text{demand_target_zone}].\text{rolling}(w).\text{mean}()$
 - 8: append demand_roll $_w$ to feat_cols
 - 9: **end for**
 - 10: drop rows with NaN (initial ≤ 168 h of each Monte Carlo year, due to lag/rolling windows) ▷ 17 total features, 26,280 $\rightarrow$ 26,112 rows
 - 11: $X \leftarrow \mathcal{D}_{\text{train}}[\text{feat_cols}]; y \leftarrow \mathcal{D}_{\text{train}}[\text{target}]$
 - 12: split (X, y) chronologically (no shuffling) into 80% train / 10% val / 10% test
 - 13: fit StandardScaler on train split only; transform all three splits
 - 14: **return** $X_{\text{tr}}, y_{\text{tr}}, X_{\text{val}}, y_{\text{val}}, X_{\text{te}}, y_{\text{te}}$, fitted scaler
-

B.6 XGBoost Surrogate Model (Section 4.4) — Gradient-Boosted Trees

Algorithm 6 TrainSurrogate — Optuna-Tuned XGBoost Price Surrogate

Require: $X_{\text{tr}}, y_{\text{tr}}, X_{\text{val}}, y_{\text{val}}, X_{\text{te}}, y_{\text{te}}$ from Algorithm B.5

Ensure: Trained XGBoost model M^* , benchmark metrics table

- 1: $\hat{y}_{\text{te}}^{\text{pers}} \leftarrow y_{\text{te}}$ shifted by 1 ▷ Baselines
 - 2: fit $\hat{y}^{\text{LR}} \leftarrow \text{LinearRegression}(X_{\text{tr}}, y_{\text{tr}})$ ▷ persistence: tomorrow's price = today's
 - 3: record MAE, RMSE, R^2 for both on X_{te} ▷ Bayesian hyperparameter search (Optuna)
 - 4: **function** OBJECTIVE(trial)
 - 5: sample $\theta = \{n_{\text{est}} \in [100, 400] \text{ step} = 100, \text{depth} \in [3, 9], \text{lr} \in [0.01, 0.15] \text{ log, sub} \in [0.7, 1.0], \text{col} \in [0.6, 1.0]\}$
 - 6: $M_{\theta} \leftarrow \text{XGBRegressor}(\theta, \text{early_stopping_rounds} = 15)$
 - 7: fit M_{θ} on $(X_{\text{tr}}, y_{\text{tr}})$ with validation set $(X_{\text{val}}, y_{\text{val}})$ for early stopping
 - 8: **return** MAE($y_{\text{val}}, M_{\theta}.\text{predict}(X_{\text{val}})$) ▷ objective to minimise
 - 9: **end function**
 - 10: $\theta^* \leftarrow \arg \min_{\theta} \text{Optuna.minimize}(\text{Objective}, n_{\text{trials}} = 25)$ ▷ 25-trial Bayesian/TPE search over θ
 - 11: $M^* \leftarrow \text{XGBRegressor}(\theta^*).\text{fit}(X_{\text{tr}} \cup X_{\text{val}}, y_{\text{tr}} \cup y_{\text{val}})$ ▷ Final fit on train $\cup$ validation, evaluate once on held-out test
 - 12: $\hat{y}_{\text{te}}^{\text{XGB}} \leftarrow M^*.\text{predict}(X_{\text{te}})$
 - 13: record MAE = 0.18 €/MWh, RMSE = 1.02 €/MWh, $R^2 = 0.9843$ ▷ reported result
 - 14: persist M^* and the fitted scaler to disk
 - 15: **return** M^* , metrics table (Table 4.6)
-

B.7 Scenario Reduction via PCA and K-Means Clustering (Section 4.5)

Algorithm 7 ReduceScenarios — PCA Dimensionality Reduction + K-Means Clustering

Require: Number of synthetic weather years $N_{\text{scenarios}}$ (20), hours per scenario (168)

Ensure: K^* representative scenarios with probability weights w_k , $\sum_k w_k = 1$

▷ Build scenario feature matrix

- 1: **for** $s = 0, 1, \dots, N_{\text{scenarios}} - 1$ **do**
- 2: generate $\text{wind}_s \in \mathbb{R}^{168}$, $\text{solar}_s \in \mathbb{R}^{168}$ via seeded sinusoidal + noise model (seed = $31s$)
- 3: $v_s \leftarrow [\text{wind}_s \parallel \text{solar}_s] \in \mathbb{R}^{336}$ ▷ concatenate wind and solar profiles
- 4: **end for**
- 5: $Y \leftarrow [v_0; v_1; \dots; v_{N_{\text{scenarios}}-1}] \in \mathbb{R}^{20 \times 336}$ ▷ PCA: retain 95% of variance
- 6: $Y_{\text{pca}} \leftarrow \text{PCA}(n_{\text{components}} = 0.95).\text{fit_transform}(Y)$ ▷ 336 $\rightarrow$ 18 dimensions, 18.7:1 compression
- ▷ K-means sweep with silhouette validation
- 7: **for** $k \in \{3, 5, 8, 10\}$ **do**
- 8: $\text{labels}_k \leftarrow \text{KMeans}(k, n_{\text{init}} = 10).\text{fit_predict}(Y_{\text{pca}})$
- 9: $\text{sil}_k \leftarrow \text{SilhouetteScore}(Y_{\text{pca}}, \text{labels}_k)$ ▷ Rousseeuw, 1987
- 10: **end for**
- 11: $K^* \leftarrow \arg \max_k \text{sil}_k$ ▷ $K^* = 8$, $\text{sil} = 0.0205$ in this run
- 12: $\text{labels} \leftarrow \text{KMeans}(K^*).\text{fit_predict}(Y_{\text{pca}})$ ▷ Extract one representative (nearest-to-centroid) scenario per cluster
- 13: **for** $c = 0, 1, \dots, K^* - 1$ **do**
- 14: $\text{mask} \leftarrow \{i : \text{labels}_i = c\}$
- 15: $w_c \leftarrow |\text{mask}| / N_{\text{scenarios}}$
- 16: $\text{rep}_c \leftarrow \arg \min_{i \in \text{mask}} \|Y_{\text{pca},i} - \text{centroid}_c\|_2$
- 17: **end for**
- 18: **return** $\{(\text{rep}_c, w_c)\}_{c=0}^{K^*-1}$

B.8 The Missing-Money Algorithm: Hybrid EVA Iteration Loop (Section 4.6)

Algorithm B.8 computes the cost side of this ratio (annualised CAPEX + FOM, via the standard capital recovery factor), and Algorithm B.9 computes the revenue side (from the LP dual price, Algorithm B.3) and adjusts capacity toward the point where revenue/cost $\rightarrow 1$.

Algorithm 8 AnnualisedCost — Capital Recovery Factor (Missing-Money Cost Side)

Require: Installed capacity p_{nom} [MW], CAPEX [€/kW], fixed O&M [€/kW/yr], WACC, asset life L [years]

Ensure: Annualised fixed cost [€/yr]

- 1: $\text{CRF} \leftarrow \frac{\text{WACC}(1+\text{WACC})^L}{(1+\text{WACC})^L - 1}$ $\triangleright$ capital recovery factor, Eq. (4.3) of thesis
 - 2: $\text{AnnualisedCost} \leftarrow (\text{CAPEX} \times 1000 \times p_{\text{nom}}) \times \text{CRF} + \text{FOM} \times 1000 \times p_{\text{nom}}$
 - 3: **return** AnnualisedCost
-

Algorithm 9 RunHybridEVA — Missing-Money Hybrid EVA Iteration Loop

Require: PyPSA network $\mathcal{N}$, trained surrogate model M^* , convergence tolerance τ ,
max iterations = 8

Ensure: Converged network $\mathcal{N}^*$, iteration history

```

1: caps  $\leftarrow \mathcal{N}.\text{generators}[p_{\text{nom}}].\text{copy}()$ 
2: history  $\leftarrow \emptyset$ 
3: for  $i = 1, 2, \dots, 8$  do
4:   use_lp  $\leftarrow (\text{surrogate unavailable}) \vee (i > 4)$ 
5:   if use_lp then
6:     res  $\leftarrow \text{RUNDISPATCH}(\mathcal{N})$   $\triangleright$  Algorithm B.3 — full LP re-solve
7:      $\lambda \leftarrow \text{res.prices}$ 
8:      $\text{rev}_g \leftarrow \sum_t \lambda_{\text{bus}(g),t} \cdot p_{g,t}$  for each generator  $g$   $\triangleright$  exact revenue from LP dual
        $\times$  primal
9:   else
10:    build feature vector from current  $\mathcal{N}$  state (loads, hour/day/month)  $\triangleright$ 
      simplified feature set — see analysis below
11:     $\lambda_{\text{DE}} \leftarrow M^*.\text{predict}(\text{features})$   $\triangleright$  surrogate replaces the LP solve entirely
12:     $\lambda_{\text{FR}} \leftarrow 0.85\lambda_{\text{DE}}; \lambda_{\text{NL}} \leftarrow 0.92\lambda_{\text{DE}}$   $\triangleright$  fixed cross-zone price ratios, not
      re-estimated
13:     $\text{rev}_g \leftarrow 0$  for all  $g$   $\triangleright$  revenue not computed in surrogate branch — see
      analysis
14:    end if
15:    for each generator  $g$  with carrier  $\in \{\text{gas, coal, nuclear, wind, solar}\}$  do
16:       $\text{ann}_g \leftarrow \text{ANNUALISED COST}(p_{\text{nom}}(g), \text{CAPEX}(g), \text{FOM}(g))$   $\triangleright$  Algorithm B.9
17:       $\text{ratio}_g \leftarrow \text{rev}_g / (\text{ann}_g + \epsilon)$   $\triangleright$  revenue-to-cost ratio  $\approx 1$  at the missing-money
      equilibrium
18:       $p_{\text{nom}}^{\text{new}}(g) \leftarrow \text{clip}(p_{\text{nom}}(g) \cdot (0.7 + 0.3 \cdot \text{ratio}_g), 0.1p_{\text{nom}}(g), 2.0p_{\text{nom}}(g))$   $\triangleright$  ratio
       $> 1$  (profitable)  $\Rightarrow$  capacity grows; ratio  $< 1$  (missing money)  $\Rightarrow$  capacity shrinks;
      damped by the 0.7/0.3 blend and clipped to  $[0.1\times, 2\times]$  to prevent divergence
19:    end for
20:    cap_change  $\leftarrow \max_g |p_{\text{nom}}^{\text{new}}(g) - p_{\text{nom}}(g)| / p_{\text{nom}}(g)$   $\triangleright$  convergence statistic
21:    history  $\leftarrow \text{history} \cup \{(i, \text{method}, \text{cap\_change}, \text{mean}(AD))\}$ 
22:    update  $\mathcal{N}.\text{generators}[p_{\text{nom}}] \leftarrow p_{\text{nom}}^{\text{new}}$ 
23:    if cap_change  $< \tau$  and use_lp then
24:      break  $\triangleright$  converged: missing-money gap closed to within tolerance
25:    end if
26: end for
27: return  $\mathcal{N}$ , history

```

Analysis of the Missing-Money Algorithm (Economic Interpretation of the Update Rule)

The line $p_{\text{nom}}^{\text{new}} = p_{\text{nom}} \cdot (0.7 + 0.3 \cdot \text{ratio})$ is a damped, first-order proportional controller driving the revenue-to-annualised-cost ratio toward 1 — the textbook zero-profit / no-missing-money condition for the marginal entrant. When $\text{ratio}_g > 1$ (the technology earns more than its annualised fixed cost, i.e. money is not missing and entry is attractive), capacity is scaled up toward $2\times$; when $\text{ratio}_g < 1$ (missing money), capacity shrinks toward $0.1\times$. The 0.7/0.3 blend is a damping factor (only 30% of the “full” adjustment implied by the ratio is applied per iteration), which plays the same numerical role as the contraction-mapping step size α_i discussed in Section 2.1.4 convergence analysis — it keeps the iteration inside the basin where the Banach fixed-point argument (and the expanded convergence discussion) is expected to apply.

Bibliography

- [1] S. Simoes et al. *Methodological Improvements of Resource Adequacy Assessments: Economic Viability Assessment (EVA)*. Tech. rep. Work Package 3 Report. European Commission / BMWK, 2021. URL: https://www.bundeswirtschaftsministerium.de/Redaktion/EN/Publikationen/Studien/methodological-improvements-of-resource-adequacy-assessments-work-package-3.pdf?__blob=publicationFile&v=1.
- [2] ENTSO-E. *European Resource Adequacy Assessment 2024 Edition - Annex 2: Methodology*. Tech. rep. ENTSO-E, 2024. URL: https://eepublicdownloads.blob.core.windows.net/public-cdn-container/clean-documents/sdc-documents/ERAA/2024/report/ERAA_2024_Annex_2_Methodology.pdf.
- [3] ENTSO-E. *European Resource Adequacy Assessment 2024 Edition - Executive Report*. Tech. rep. ENTSO-E, 2024. URL: https://eepublicdownloads.blob.core.windows.net/public-cdn-container/clean-documents/sdc-documents/ERAA/2024/report/ERAA_2024_Executive_Report.pdf.
- [4] Jonas Hörsch et al. “PyPSA-Eur: An Open Optimisation Model of the European Transmission System”. In: *Energy Strategy Reviews* 22 (2018), pp. 207–215. DOI: [10.1016/j.esr.2018.08.012](https://arxiv.org/abs/1806.01613). URL: <https://arxiv.org/abs/1806.01613>.
- [5] European Parliament and Council of the European Union. *Regulation (EU) 2019/943 of the European Parliament and of the Council of 5 June 2019 on the internal market for electricity (recast)*. OJ L 158, 14.6.2019, pp. 54–124. 2019. URL: <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32019R0943>.
- [6] ENTSO-E. *European Resource Adequacy Assessment 2025 Edition - Annex 2: Methodology*. Tech. rep. ENTSO-E, 2025. URL: https://eepublicdownloads.blob.core.windows.net/public-cdn-container/clean-documents/sdc-documents/ERAA/ERAA_2025/ERAA%202025%20Annex%202_Amended%20edition%20following%20ACER%27s%20decision.pdf.
- [7] IEEE Working Group on Resource Adequacy. “Electric Grid Reliability in an Era of Unprecedented Uncertainty: A Review of Advances in Electric Sector Resource Adequacy Assessment and Planning”. In: *IEEE Transactions on Power Systems* (2025). URL: <https://ieeexplore.ieee.org/document/11278497>.

- [8] ACER. *Decision on the European Resource Adequacy Assessment 2024*. Tech. rep. Agency for the Cooperation of Energy Regulators, 2025. URL: <https://www.acer.europa.eu/sites/default/files/documents/Individual%20Decisions/ACER-Decision-07-2025-ERAA-2024.pdf>.
- [9] ENTSO-E. *European Resource Adequacy Assessment 2024 Edition - Annex 4: Case Study – Application of Revenue-Based EVA Approach on European Scale*. Tech. rep. ENTSO-E, 2024. URL: https://eepublicdownloads.blob.core.windows.net/public-cdn-container/clean-documents/sdc-documents/ERAA/2024/report/ERAA_2024_Annex_4_Case_Study_Revenue-Based_EVA.pdf.
- [10] ENTSO-E. *European Resource Adequacy Assessment 2025 Edition - Annex 5: Proof of Concept: Application of Revenue-Based EVA Approach on European Scale*. Tech. rep. ENTSO-E, 2025. URL: https://eepublicdownloads.blob.core.windows.net/public-cdn-container/clean-documents/sdc-documents/ERAA/ERAA_2025/ERAA%202025%20Annex%205_Amended%20edition%20following%20ACER%27s%20decision.pdf.
- [11] Stefan Banach. “Sur les opérations dans les ensembles abstraits et leur application aux équations intégrales”. In: *Fundamenta Mathematicae* 3 (1 1922). <https://eudml.org/doc/213289>, pp. 133–181. ISSN: 0016-2736. URL: <https://matwbn.icm.edu.pl/ksiazki/fm/fm3/fm3120.pdf>.
- [12] RTE International. *Open-Source Solution for European Data Collection: PEM-MDB Deployment*. Tech. rep. PEMMDB centralises and certifies TSO data quality for ENTSO-E studies. RTE International, 2026. URL: <https://www.rte-international.com/en/open-source-solution-entsoe-data-collection/>.
- [13] ENTSO-E. *European Resource Adequacy Assessment 2025: Annex Ib - Input Data and Assumptions*. Tech. rep. ACER Decision No 06-2026. European Network of Transmission System Operators for Electricity, 2026. URL: https://www.acer.europa.eu/sites/default/files/documents/Individual%20Decisions_annex/ACER-Decision-06-2026-ERAA-2025-Annex-Ib-Input-data-assumptions.pdf.
- [14] ENTSO-E. *European Resource Adequacy Assessment 2023: Annex 2 - Methodology*. Tech. rep. Discusses PEMMDB, PECD, and data collection processes. European Network of Transmission System Operators for Electricity, 2023. URL: https://www.entsoe.eu/outlooks/eraa/2023/report/ERAA_2023_Annex_2_Methodology.pdf.

- [15] Copernicus Climate Change Service (C3S). *Pan-European Climate Database Version 4.2: Data Description and Improvements*. Tech. rep. PECD v4.2 used for ERAA 2026, includes 1950-2024 historical data and projections to 2100. Copernicus Climate Change Service, 2026. URL: <https://energy.hub.copernicus.eu/pan-european-climate-database-update-what-it-means-users-energy-sector>.
- [16] N. Gonzalez-Reviriego et al. *Advancing C3S Energy Services in support to Europe's energy transition: co-development of the Pan-European Climate Database with ENTSO-E*. EGU25-19306, describes PECD co-development between C3S and ENTSO-E. 2025. DOI: [10.5194/egusphere-egu25-19306](https://doi.org/10.5194/egusphere-egu25-19306).
- [17] ENTSO-E. *Demand Forecasting Methodology*. Tech. rep. Describes the Demand Forecasting Toolbox (DFT) and its integration in ERAA. European Network of Transmission System Operators for Electricity, 2019. URL: https://eepublicdownloads.entsoe.eu/clean-documents/sdc-documents/MAF/2020/Demand_forecasting_methodology_V1_1.pdf.
- [18] George B. Dantzig. *Linear Programming and Extensions*. Princeton, NJ: Princeton University Press, 1963. ISBN: 978-0-691-05913-6. URL: <https://press.princeton.edu/books/paperback/9780691059136/linear-programming-and-extensions>.
- [19] F. C. Schweppe et al. *Spot Pricing of Electricity*. New York, NY: Springer, 1988. ISBN: 978-1-4612-8922-7. DOI: [10.1007/978-1-4613-1683-1](https://doi.org/10.1007/978-1-4613-1683-1). URL: <https://link.springer.com/book/10.1007/978-1-4613-1683-1>.
- [20] Tianqi Chen and Carlos Guestrin. “XGBoost: A Scalable Tree Boosting System”. In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. The seminal paper introducing the XGBoost algorithm. 2016, pp. 785–794. DOI: [10.1145/2939672.2939785](https://doi.org/10.1145/2939672.2939785).
- [21] C. Albertyn. “A Comparative Study of Gradient Boosting Algorithms and Transformer Neural Networks for Classification Problems in the Tabular Data Domain”. XGBoost outperformed both TabTransformer and FT-Transformer in tabular data classification. MA thesis. Stellenbosch University, 2025. URL: <https://scholar.sun.ac.za/items/1426022c-fbf0-44eb-ab74-43df924264d6>.
- [22] Keyong Hu, Qingqing Yang, Lei Lu, et al. “Two-Stage Distributionally Robust Optimal Scheduling for Integrated Energy Systems Considering Uncertainties in Renewable Generation and Loads”. In: *Mathematics* 13.9 (2025). Discusses K-means clustering for scenario reduction in energy systems, p. 1439. DOI: [10.3390/math13091439](https://doi.org/10.3390/math13091439).

- [23] Ali Almohammed and Mohammed Abido. “RES on Power Operation: K-means Clustering over Elbow Approach”. In: *WSEAS Transactions on Power Systems* 15 (2020). Combines PCA dimensionality reduction with K-means clustering for power system operation modes, pp. 214–221. URL: <https://pure.kfupm.edu.sa/en/publications/res-on-power-operation-k-means-clustering-over-elbow-approach/>.
- [24] Krishna Mohan et al. “Scenario Generation and Reduction Based Uncertainty Mitigation in a Microgrid”. In: *IEEE 4th International Conference on Sustainable Energy and Future Electric Transportation (SEFET)*. 2024, pp. 1–5. DOI: [10.1109/SEFET61574.2024.10718010](https://doi.org/10.1109/SEFET61574.2024.10718010).
- [25] Yajie Tang et al. “Selection of Typical Operating Scenarios for High-Penetration Renewable Energy Power Systems Based on Clustering and Probabilistic Modeling”. In: *IEEE Sustainable Power and Energy Conference (iSPEC)*. 2025. URL: <https://ieeexplore.ieee.org/document/11239740>.